

HW03

Minyoung Do

2/4/2020

Egalitarianism and income

```
gss_test <- read_csv("gss_test.csv")
gss_train <- read_csv("gss_train.csv")
```

(20 points) Perform polynomial regression to predict `egalit_scale` as a function of `income06`. Use and plot 10-fold cross-validation to select the optimal degree d for the polynomial based on the MSE. Plot the resulting polynomial fit to the data, and also graph the average marginal effect (AME) of `income06` across its potential values. Be sure to provide substantive interpretation of the results.

referenced this post on Stack Overflow: <https://stackoverflow.com/questions/43686539/using-cv>

```
k <- 10
fold <- sample(k, nrow(gss_train), replace = TRUE)

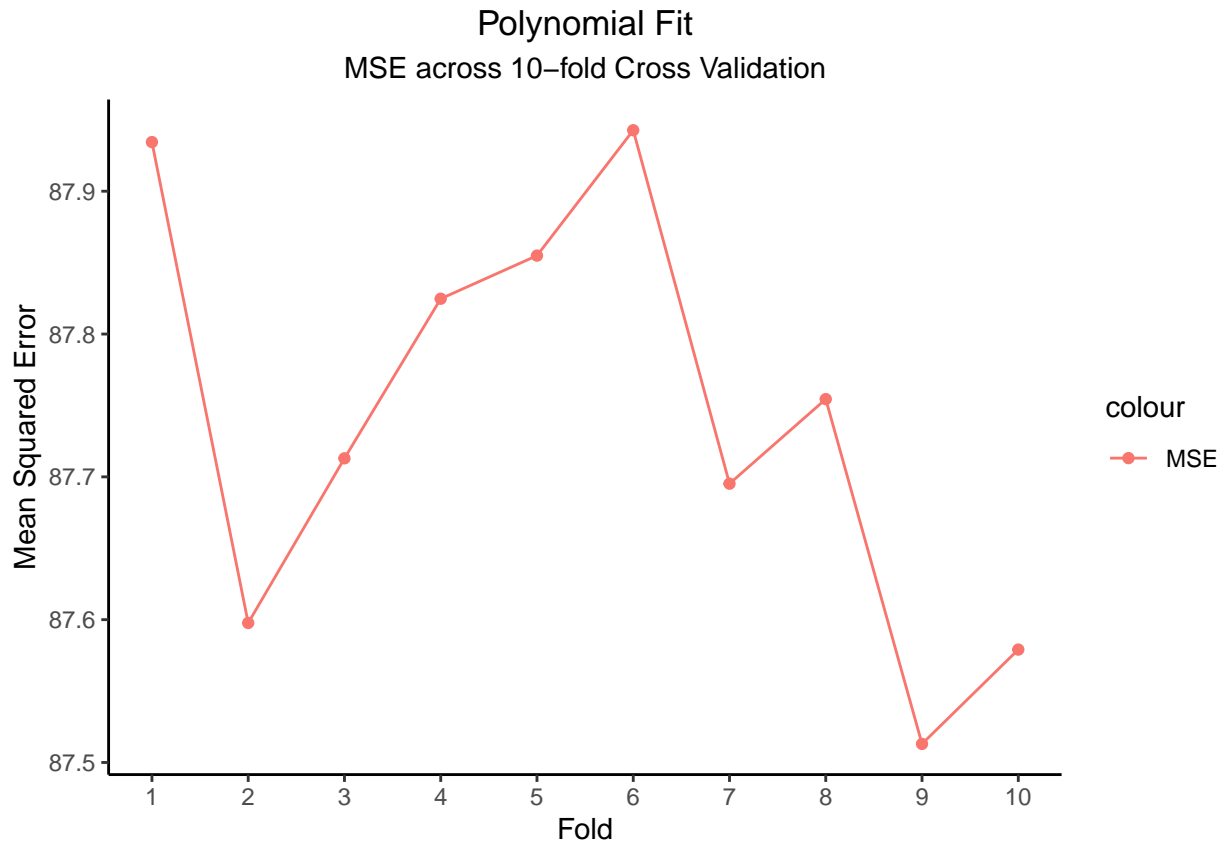
## For each span from 1 to 10 we can calculate the CV test error:
mse <- numeric(k)
span <- seq(1, 10, by = 1)
cv <- numeric(length(span))

for (j in seq_along(span))
{
  for (i in seq_len(k))
  {
    take <- fold == i
    foldi <- gss_train[take, ]
    foldOther <- gss_train[!take, ]
    f <- glm(egalit_scale ~ poly(income06, span[j]), data=foldOther)
    pred <- predict(f, foldi)
    mse[i] <- mean((pred - foldi$egalit_scale)^2, na.rm = TRUE)
  }
  cv[j] <- mean(mse)
  result <- cbind(span, cv) %>%
    as.tibble() %>%
    rename(Fold = "span", MeanMSE = "cv")
}
```

```
result %>%
  kable() %>%
  kable_styling(full_width = F)
```

Fold	MeanMSE
1	87.93442
2	87.59767
3	87.71297
4	87.82472
5	87.85491
6	87.94272
7	87.69522
8	87.75443
9	87.51300
10	87.57905

```
ggplot(result, aes(Fold, MeanMSE, color = "MSE")) +
  geom_point() +
  geom_line() +
  theme_classic() +
  theme(plot.title = element_text(hjust = 0.5), plot.subtitle = element_text(hjust = 0.5)) +
  scale_x_continuous(breaks=c(1:10)) +
  labs(title = "Polynomial Fit",
       subtitle = "MSE across 10-fold Cross Validation",
       x = "Fold",
       y = "Mean Squared Error")
```




As seen in the graph, the second order polynomial fits the data best with the mean MSE value of 88.01.

```
poly_ame <- lm(egalit_scale ~ income06 + I(income06^2), data = gss_train)
margins(poly_ame)
```

```
## income06
## -0.4303
```

```
cplot(object = poly_ame, x = "income06", what = "effect")
```



This graph shows the marginal effect of `income06` in a negative line, which indicates that the marginal effect of `income06` decreases as the income level increases. The average marginal effect is -0.43, which corresponds with the graph with a negative slope.

(20 points) Fit a step function to predict `egalit_scale` as a function of `income06`, and perform 10-fold cross-validation to choose the optimal number of cuts. Plot the fit and interpret the results.

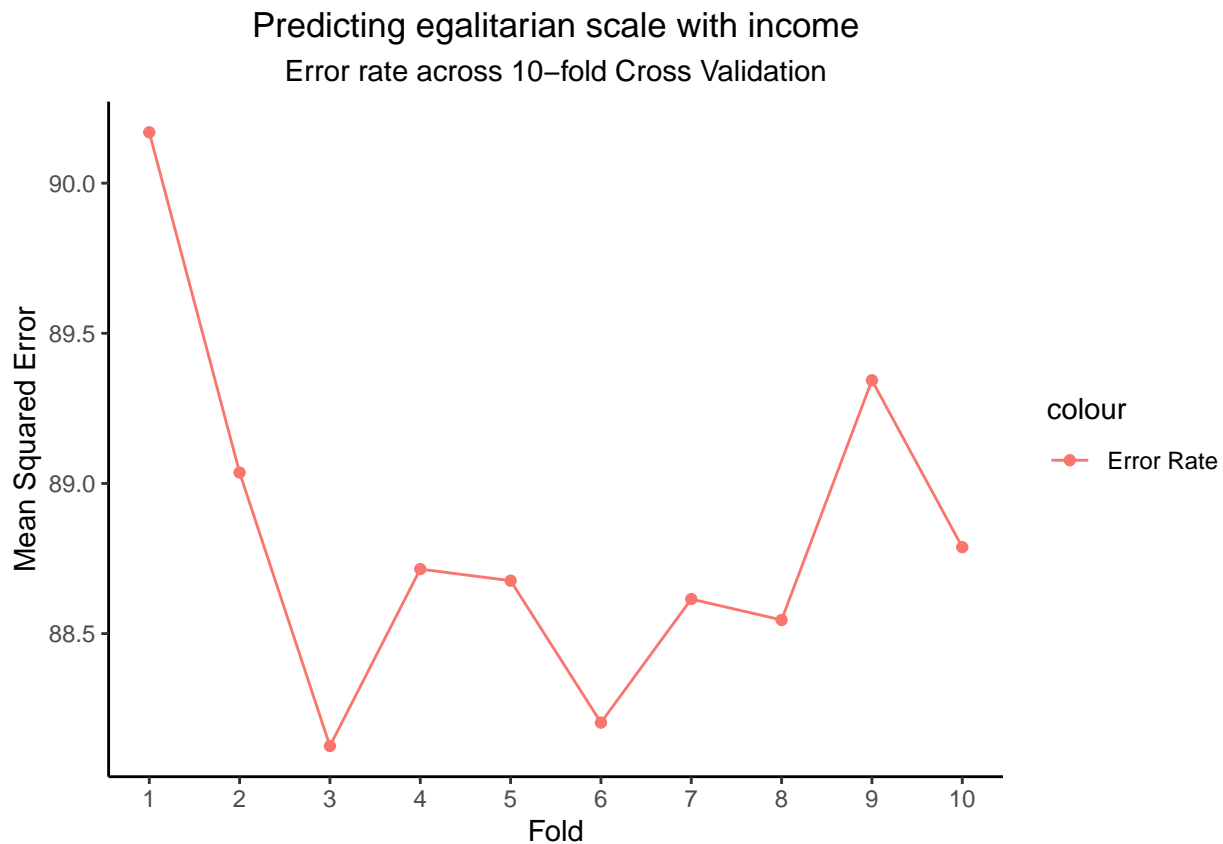
```
# referenced this Stack Overflow post: https://stackoverflow.com/questions/42190337/cross-validation-for-step-function

step_fun <- vector(mode = "numeric", length = 10)

for (i in 1:10) {
  cuts <- levels(cut(gss_train$income06, i + 1))
  breaks <- unique(c(as.numeric(sub("\\((.+),.*", "\\1", cuts)),
                     as.numeric(sub("[^,]*,([^\]]*)\\]", "\\1", cuts))))
  interval <- glm(egalit_scale ~ cut(income06, unique(breaks)), data = gss_train)
  step_fun[i] <- cv.glm(gss_train, interval, K = 10)$delta[1]
  error_result <- step_fun %>%
    as.tibble() %>%
    mutate(id = seq_len(n()))
}

error_result %>%
  ggplot() +
  geom_point(aes(id, value, color = "Error Rate")) +
  geom_line(aes(id, value, color = "Error Rate")) +
```

```
theme_classic() +
theme(plot.title = element_text(hjust = 0.5), plot.subtitle = element_text(hjust = 0.5)) +
scale_x_continuous(breaks=c(1:10)) +
labs(title = "Predicting egalitarian scale with income",
      subtitle = "Error rate across 10-fold Cross Validation",
      x = "Fold",
      y = "Mean Squared Error")
```



```
min(error_result[,1])
```

```
## [1] 88.12559
```

The graph suggests that the 6 cuts are the optimal number for the fit, with the error rate of 88.28 out of 10 models. Let's take a closer look at this model:

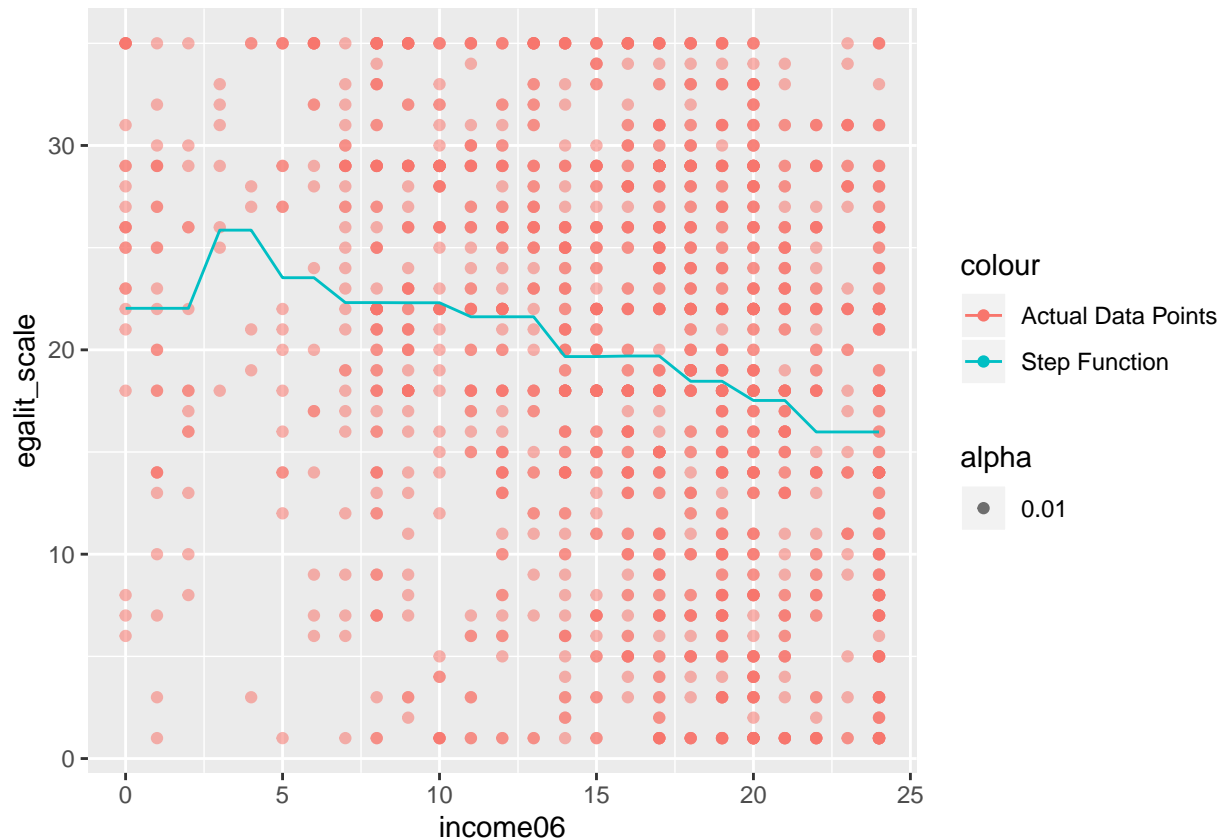
```
step_cut <- levels(cut(gss_train$income06, i + 1))
step_breaks <- unique(c(as.numeric(sub("\\((.+),.*", "\\1", cuts)),
                        as.numeric(sub("[^,]*,([^\]]*)\\]", "\\1", cuts))))
step_fit <- glm(egalit_scale ~ cut(income06, unique(breaks)), data = gss_train)
step_pred <- predict(step_fit, gss_train)
```

```

step_df <- tibble(step_pred, gss_train$income06) %>%
  rename(income = "gss_train$income06")

ggplot() +
  geom_point(data = gss_train, aes(income06, egalit_scale, alpha = 0.01, color = "Actual Data Points"))
  geom_line(data = step_df, aes(income, step_pred, color = "Step Function"))

```



The line graph shows how many bins there are for the `income06` variable. There are 9 bins total, and there is a decreasing pattern overall as the income increases.

(20 points) Fit a natural regression spline to predict `egalit_scale` as a function of `income06`. Use 10-fold cross-validation to select the optimal number of degrees of freedom, and present the results of the optimal model.

```

# function to simplify things
egalit_spline <- function(splits, df = NULL){
  # estimate the model on each fold
  model <- glm(egalit_scale ~ ns(income06, df = df),
    data = analysis(splits))

  model_acc <- augment(model, newdata = assessment(splits)) %>%
    mse(egalit_scale, estimate = .fitted)
}

```

```

  mean(model_acc$.estimate)
}


knots_fun <- function(splits, knots){
  egalit_spline(splits, df = knots + 3)
}

# estimate CV error for knots in 0:10
results <- vfold_cv(gss_train, v = 10)

tidyr::expand(results, id, knots = 1:10) %>%
  left_join(results) %>%
  mutate(acc = map(splits, knots, tune_over_knots)) %>%
  group_by(knots) %>%
  summarize(acc = mean(acc)) %>%
  mutate(err = 1 - acc) %>%
  ggplot(aes(knots, err)) +
  geom_point() +
  geom_line() +
  scale_y_continuous(labels = scales::percent) +
  labs(title = "Optimal number of knots for natural cubic spline regression",
       x = "Knots",
       y = "10-fold CV error")

```

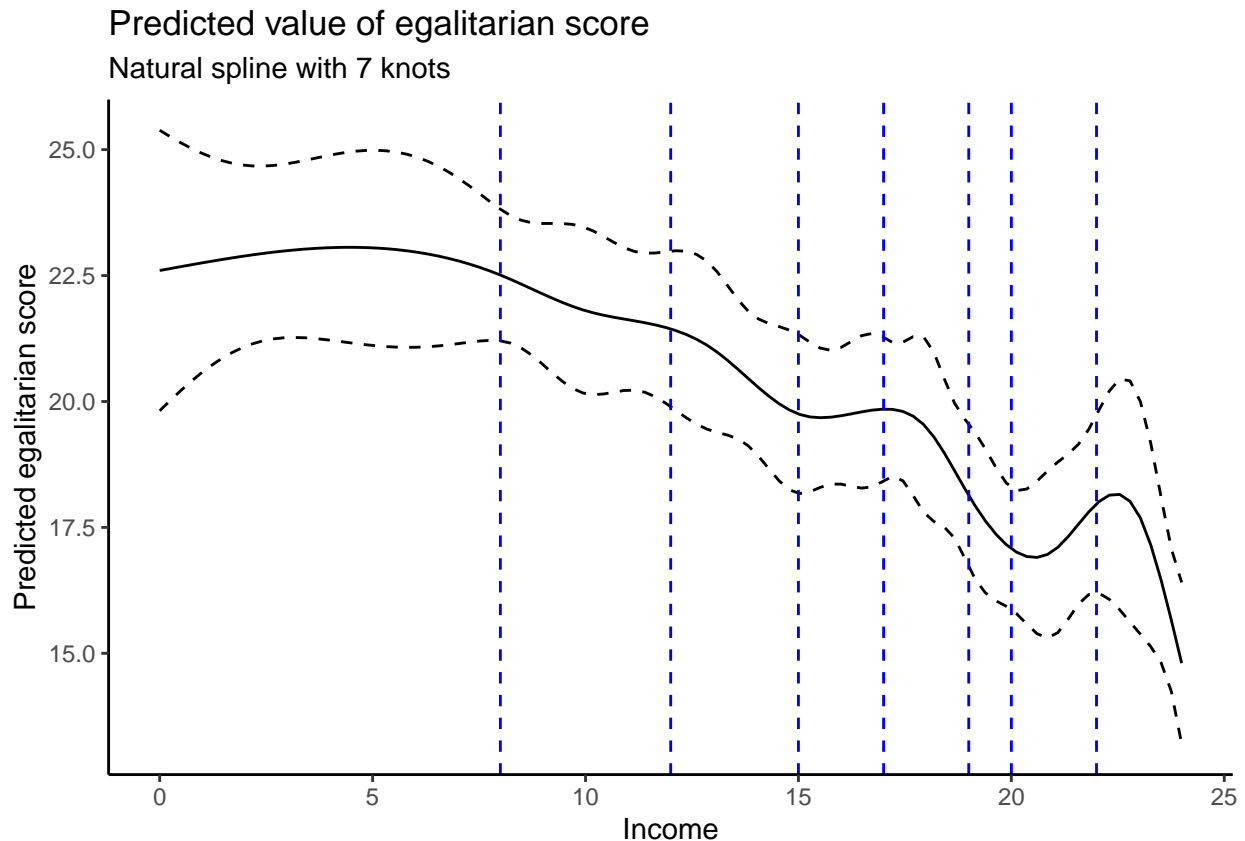
Optimal number of knots for natural cubic spline regression



```
df_glm <- glm(egalit_scale ~ ns(income06, df = 10), data = gss_train) %>%
  cplot("income06", what = "prediction", n = 100, draw = FALSE) %>%
  as.tibble()
```

```
##      xvals    yvals    upper    lower
## 1  0.000000  22.60034  25.38579  19.81488
## 2  0.242424  22.63784  25.25913  20.01656
## 3  0.484848  22.67513  25.14081  20.20944
## 4  0.727272  22.71197  25.03245  20.39149
## 5  0.969697  22.74814  24.93568  20.56061
## 6  1.212121  22.78343  24.85218  20.71468
## 7  1.454545  22.81761  24.78349  20.85173
## 8  1.696969  22.85046  24.73086  20.97006
## 9  1.939393  22.88175  24.69505  21.06845
## 10 2.181818  22.91127  24.67619  21.14636
## 11 2.424242  22.93879  24.67354  21.20405
## 12 2.666667  22.96410  24.68556  21.24264
## 13 2.909090  22.98696  24.70994  21.26399
## 14 3.151515  23.00717  24.74380  21.27054
## 15 3.393939  23.02448  24.78395  21.26502
## 16 3.636363  23.03870  24.82712  21.25027
## 17 3.878787  23.04958  24.87012  21.22904
## 18 4.121212  23.05692  24.90995  21.20388
## 19 4.363636  23.06048  24.94389  21.17707
## 20 4.606060  23.06005  24.96950  21.15060
```

```
df_glm %>%
  ggplot(aes(x = xvals)) +
  geom_line(aes(y = yvals)) +
  geom_line(aes(y = upper), linetype = 2) +
  geom_line(aes(y = lower), linetype = 2) +
  theme_classic() +
  geom_vline(xintercept = attr(bs(gss_train$income06, df = 10), "knots"),
             linetype = 2, color = "blue") +
  labs(title = "Predicted value of egalitarian score",
       subtitle = "Natural spline with 7 knots",
       x = "Income",
       y = "Predicted egalitarian score")
```

We can observe a similar pattern found in the polynomial regression and step-function-applied model, i.e., the decreasing pattern of egalitarian scale with increasing income levels. This natural regression model has 7 knots total, whereas the step function model has 9 bins. Despite the smaller number of knots compared to 9 bins, this regression line represents the data better with a smoother line.

Egalitarianism and everything

(20 points total) Estimate the following models using all the available predictors (be sure to perform appropriate data pre-processing (e.g., feature standardization) and hyperparameter tuning (e.g. lambda for PCR/PLS, lambda and alpha for elastic net). Also use 10-fold cross-validation for each model to estimate the model's performance using MSE):

a. (5 points) Linear regression

```
ten_fold <- trainControl(method = "cv", number = 10)

gss_lm <- train(egalit_scale ~ ., data = gss_train,
               method = "lm", trControl = ten_fold)
```

```
pred_train <- predict(gss_lm, gss_test)
Metrics::mse(gss_test$egalit_scale, pred_train)
```

```
## [1] 63.21363
```

b. (5 points) Elastic net regression

```
gss_train_x <- model.matrix(egalit_scale ~ ., gss_train)[, -1]
gss_train_y <- gss_train$egalit_scale
```

```
gss_test_x <- model.matrix(egalit_scale ~ ., gss_test)[, -1]
gss_test_y <- gss_test$egalit_scale
```

Now, more efficient grid search for varying alpha

```
fold_id <- sample(1:10, size = length(gss_train_y), replace = TRUE) # maintain the same folds
```

search across a range of alphas

```
tuning_grid <- tibble::tibble(
  alpha      = seq(0, 1, by = .1),
  mse_min    = NA,
  mse_1se    = NA,
  lambda_min = NA,
  lambda_1se = NA
)
```

```
for(i in seq_along(tuning_grid$alpha)) {
  # fit CV model for each alpha value
  fit <- cv.glmnet(gss_train_x,
                  gss_train_y,
                  alpha = tuning_grid$alpha[i],
                  foldid = fold_id)
```

extract MSE and lambda values

```
tuning_grid$mse_min[i] <- fit$cvm[fit$lambda == fit$lambda.min]
tuning_grid$mse_1se[i] <- fit$cvm[fit$lambda == fit$lambda.1se]
tuning_grid$lambda_min[i] <- fit$lambda.min
tuning_grid$lambda_1se[i] <- fit$lambda.1se
}
```

minimum MSE

```
min(tuning_grid[,2])
```

```
## [1] 60.09246
```

```
# minimum lambda
min(tuning_grid[,4])
```

```
## [1] 0.2127845
```


```
# combination of alpha = 1 & lambda = 0.2127845
coef_en <- coef(fit, alpha = 1)
```

c. (5 points) Principal component regression


```
# running pcr
gss_pcr <- pcr(egalit_scale ~ .,
              # just grabbing numeric features
              data = select_if(gss_train, is.numeric),
              center = TRUE,
              validation = "CV",
              segments = 10)

# extracting gss_exp, mse, pc, and cum_exp
gss_pcr_stats <- tibble(
  gss_exp = loadings(gss_pcr) %>%
    attr("explvar"),
  mse = as.vector(MSEP(gss_pcr, estimate = "CV", intercept = FALSE)$val)
) %>%
  mutate(pc = row_number(),
         cum_exp = cumsum(gss_exp) / 100)

# percent of variance explained by each model
ggplot(gss_pcr_stats, aes(pc, cum_exp)) +
  geom_line() +
  scale_y_continuous(labels = scales::percent) +
  theme_classic() +
  labs(x = expression(M),
       y = "Cumulative percent of variance explained")
```



```
# MSE
ggplot(gss_pcr_stats, aes(pc, mse)) +
  geom_line() +
  geom_vline(xintercept = which.min(gss_pcr_stats$mse), linetype = 2) +
  theme_classic() +
  labs(x = expression(M),
       y = "CV estimate of MSE")
```



This tuning process tells us that the MSE and proportion of variance doesn't improve anymore after the number of principal components passes a certain point: the optimal number is 40.


d. (5 points) Partial least squares regression

```
# running pls
gss_pls <- plsr(egalit_scale ~ .,
               data = select_if(gss_train, is.numeric),
               center = TRUE,
               validation = "CV",
               segments = 10)

# extracting numbers we need
gss_pls_stats <- tibble(
  pct_exp = loadings(gss_pls) %>% attr("explvar"),
  mse = as.vector(MSEP(gss_pls, estimate = "CV", intercept = FALSE)$val)
) %>%
  mutate(pc = row_number(),
         cum_exp = cumsum(pct_exp) / 100)
```

(20 points) For each final tuned version of each model fit, evaluate feature importance by generating feature interaction plots. Upon visual presentation, be sure to discuss the substantive results for these models and in comparison to each other (e.g., talk about feature importance, conditional effects, how these are ranked differently across different models, etc.).

```
gss_stats <- bind_rows(PCR = gss_pcr_stats,  
                      PLS = gss_pls_stats,  
                      .id = "model")  
  
# percent of variance explained by M  
ggplot(gss_stats, aes(pc, cum_exp, color = model)) +  
  geom_line() +  
  scale_y_continuous(labels = scales::percent) +  
  scale_color_brewer(type = "qual") +  
  theme_classic() +  
  labs(x = expression(M),  
       y = "Cumulative percent of variance explained",  
       color = NULL) +  
  theme(legend.position = "bottom")
```




```

# PCR explains more variance with less coefficients

# MSE comparison
ggplot(gss_stats, aes(pc, mse, color = model)) +
  geom_line() +
  geom_vline(data = tribble(
    ~pc, ~model,
    which.min(gss_pls_stats$mse), "PLS",
    which.min(gss_pcr_stats$mse), "PCR"
  ),
  aes(xintercept = pc, color = model), linetype = 2, show.legend = FALSE) +
  scale_color_brewer(type = "qual") +
  theme_classic() +
  labs(x = expression(M),
       y = "CV estimate of MSE",
       color = NULL) +
  theme(legend.position = "bottom")

```



The graph above compares the last two models, PCR and PLS. They both perform similarly, but in between 0 and 20, PCR explains the variance of the data a little better, reaching the maximum percent of variance a little before PLS. However, in the MSE plot, PLS is slightly better as its best fit is at 20, while PCR has the best fit when $M = 40$. Both simple linear regression and elastic net has way higher MSE values. In this analysis, PCR and PLS have a better prediction accuracy of the data than linear regression and elastic net.