

Homework 4: Moving Beyond Linearity

Luying Jiang

```
gss_train <- read.csv('~/Desktop/problem-set-4-master/data/gss_train.csv')
gss_test  <- read.csv('~/Desktop/problem-set-4-master/data/gss_test.csv')
```

1. Perform polynomial regression to predict `egalit_scale` as a function of `income06`. Use and plot 10-fold cross-validation to select the optimal degree `d` for the polynomial based on the MSE. Plot the resulting polynomial fit to the data, and also graph the average marginal effect (AME) of `income06` across its potential values. Be sure to provide substantive interpretation of the results.

```
library(ISLR)
library(boot)
library(caret)
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'lattice'
```

```
## The following object is masked from 'package:boot':
```

```
##
```

```
##      melanoma
```

```
## Loading required package: ggplot2
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse
```

```
## v tibble  2.1.3      v dplyr    0.8.3
```

```
## v tidyr   1.0.0      v stringr 1.4.0
```

```
## v readr   1.3.1      v forcats 0.4.0
```

```
## v purrr   0.3.3
```

```
## -- Conflicts ----- tidyverse
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## x purrr::lift()    masks caret::lift()
```

```
library(broom)
```

```
library(ggthemes)
```

```
library(knitr)
```

```
library(kableExtra)
```

```
##
```

```
## Attaching package: 'kableExtra'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      group_rows
```

```


set.seed(123)
cv.MSE <- NA
for (i in 1:10) {
  glm.fit <- glm(egalit_scale ~ poly(income06, i), data = gss_train)
  cv.MSE[i] <- cv.glm(gss_train, glm.fit, K = 10)$delta[1]
}
cv.MSE

## [1] 87.48125 87.15967 87.33259 87.60095 87.53414 87.65479 87.53396 87.20507
## [9] 87.20160 87.27881

plot( x = 1:10, y = cv.MSE, xlab = "Degree of Polynomial", ylab = "Mean Squared Error",
      type = "b", lwd = 2, pch = 19, bty = "n",
      ylim = c( min(cv.MSE) - sd(cv.MSE), max(cv.MSE) + sd(cv.MSE) ) )
# horizontal line for 1se to less complexity
abline(h = min(cv.MSE) + sd(cv.MSE) , lty = "dotted")

# where is the minimum
points( x = which.min(cv.MSE), y = min(cv.MSE), col = "red", pch = "X", cex = 1.5 )

```



optimal degree for the polynomial regression is 2.

```

library(margins)
best_poly <- lm(egalit_scale ~ income06 + I(income06^2), data = gss_train)
margin <- margins(best_poly)
margin

## Average marginal effects

## lm(formula = egalit_scale ~ income06 + I(income06^2), data = gss_train)

## income06
## -0.4507


plot_best_poly <- cplot(best_poly, "income06", draw=FALSE)

```

##	xvals	yvals	upper	lower
## 1	1	23.11631	25.28371	20.94890
## 2	2	23.15180	25.04001	21.26359
## 3	3	23.15524	24.79232	21.51817
## 4	4	23.12665	24.54195	21.71134
## 5	5	23.06600	24.29036	21.84165
## 6	6	22.97331	24.03899	21.90764
## 7	7	22.84858	23.78870	21.90846
## 8	8	22.69180	23.53894	21.84467
## 9	9	22.50298	23.28678	21.71917
## 10	10	22.28211	23.02670	21.53752
## 11	11	22.02920	22.75132	21.30708
## 12	12	21.74424	22.45297	21.03552
## 13	13	21.42724	22.12502	20.72946
## 14	14	21.07819	21.76262	20.39376
## 15	15	20.69710	21.36290	20.03130
## 16	16	20.28396	20.92501	19.64291
## 17	17	19.83878	20.45044	19.22712
## 18	18	19.36155	19.94356	18.77955
## 19	19	18.85228	19.41247	18.29210
## 20	20	18.31096	18.86919	17.75274

```
ggplot(plot_best_poly, aes(x = xvals)) +
  geom_line(aes(y = yvals), color = 'red') +
  geom_line(aes(y = upper), linetype = 2) +
  geom_line(aes(y = lower), linetype = 2) +
  labs(x = "Income", y = "Egalitarianism",
        title = "Plot of Best Fit Polynomial") +
  theme(axis.text = element_text(size = 10),
        axis.title.x = element_text(size = 10),
        axis.title.y = element_text(size = 10))
```

Plot of Best Fit Polynomial




```
plot_best_poly <- cplot(best_poly, "income06", what = 'effect', draw=FALSE)
plot_best_poly
```

##	xvals	yvals	upper	lower	factor
##	1.0000	0.0515	0.3759	-0.2729	income06
##	2.0000	0.0195	0.3216	-0.2827	income06
##	3.0000	-0.0126	0.2674	-0.2926	income06
##	4.0000	-0.0446	0.2134	-0.3027	income06
##	5.0000	-0.0767	0.1596	-0.3129	income06
##	6.0000	-0.1087	0.1061	-0.3235	income06
##	7.0000	-0.1408	0.0529	-0.3344	income06
##	8.0000	-0.1728	0.0002	-0.3458	income06
##	9.0000	-0.2048	-0.0519	-0.3578	income06
##	10.0000	-0.2369	-0.1029	-0.3708	income06
##	11.0000	-0.2689	-0.1526	-0.3853	income06
##	12.0000	-0.3010	-0.2000	-0.4020	income06
##	13.0000	-0.3330	-0.2440	-0.4221	income06
##	14.0000	-0.3651	-0.2831	-0.4470	income06
##	15.0000	-0.3971	-0.3162	-0.4781	income06
##	16.0000	-0.4292	-0.3428	-0.5155	income06
##	17.0000	-0.4612	-0.3642	-0.5583	income06
##	18.0000	-0.4932	-0.3817	-0.6048	income06
##	19.0000	-0.5253	-0.3967	-0.6538	income06
##	20.0000	-0.5573	-0.4101	-0.7045	income06
##	21.0000	-0.5894	-0.4224	-0.7563	income06
##	22.0000	-0.6214	-0.4340	-0.8089	income06
##	23.0000	-0.6535	-0.4450	-0.8619	income06

```
## 24.0000 -0.6855 -0.4557 -0.9154 income06
## 25.0000 -0.7176 -0.4660 -0.9691 income06

ggplot(plot_best_poly, aes(x = xvals)) +
  geom_line(aes(y = yvals), color = 'red') +
  geom_line(aes(y = upper), linetype = 2) +
  geom_line(aes(y = lower), linetype = 2) +
  labs(x = "Income", y = "Egalitarianism",
       title = "Plot of Best Fit Polynomial") +
  theme(axis.text = element_text(size = 10),
        axis.title.x = element_text(size = 10),
        axis.title.y = element_text(size = 10))
```

Plot of Best Fit Polynomial



The marginal effect of income level on egalitarianism is largely negative over different income levels, which indicates that as individuals make more money, they are less likely to have egalitarian preferences. The AME for income06 is -0.4507319.

2. Fit a step function to predict egalit_scale as a function of income06, and perform 10-fold cross-validation to choose the optimal number of cuts. Plot the fit and interpret the results.

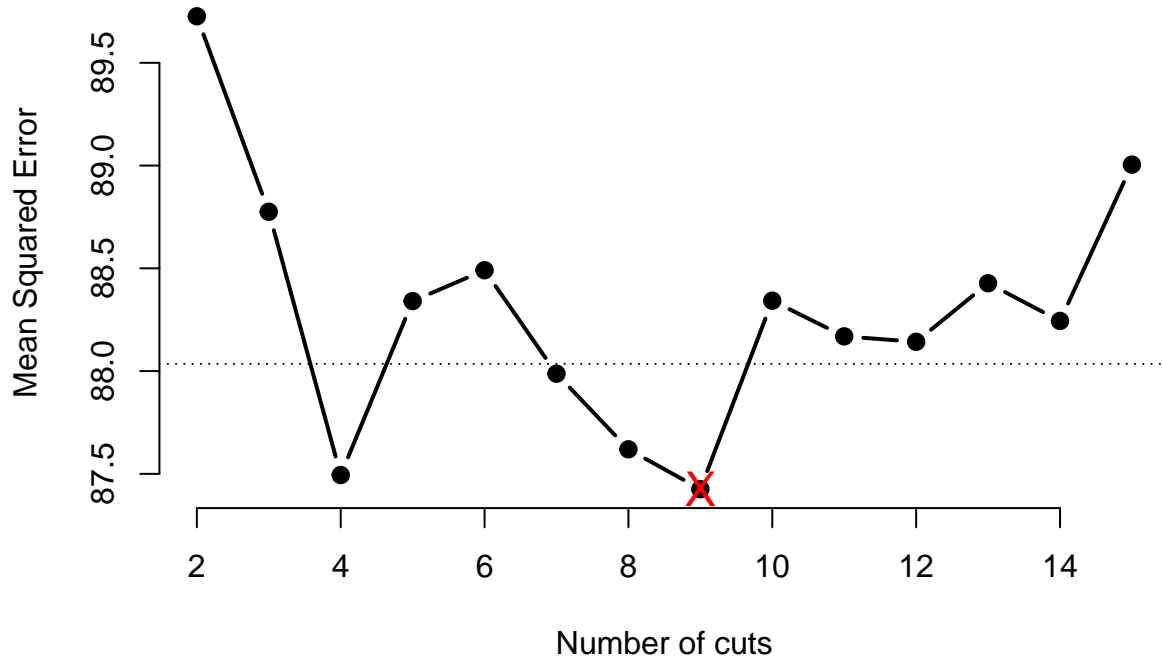
```
cv.MSE <- NA
# for each cut perform 10-fold cross-validation
for (i in 2:15) {
  gss_train$income06.cut <- cut(gss_train$income06, i)
  gss_test$income06.cut <- cut(gss_test$income06, i)
  slm.fit <- glm(egalit_scale ~ income06.cut, data = gss_train)
  cv.MSE[i] <- cv.glm(gss_train, slm.fit, K = 10)$delta[1]
}
```

```

# the first element of cv.MSE is NA because we started our loop at 2
plot(2:15, cv.MSE[-1], xlab = "Number of cuts", ylab = "Mean Squared Error",
     type = "b", pch = 19, lwd = 2, bty = "n")
abline(h = min(cv.MSE, na.rm = TRUE) + sd(cv.MSE, na.rm = TRUE) , lty = "dotted")

# highlight minimum
points( x = which.min(cv.MSE), y = min(cv.MSE, na.rm = TRUE), col = "red", pch = "X", cex = 1.5 )

```




From

the above graph, the optimal number of steps is 4. (optimal number of bins = 4)

```

lm.fit <- glm(egalit_scale ~ cut(income06, 4), data = gss_train)
incomelims <- range(gss_train$income06)
income.grid <- seq(from = incomelims[1], to = incomelims[2])
lm.pred <- predict(lm.fit, data.frame(income06 = income.grid), se = TRUE)
plot(egalit_scale ~ income06, data = gss_train, col = "darkgrey", bty = "n")
lines(income.grid, lm.pred$fit, col = "red", lwd = 2)
matlines(income.grid, cbind( lm.pred$fit + 2* lm.pred$se.fit,
                             lm.pred$fit - 2* lm.pred$se.fit),
         col = "red", lty = "dashed")

```



There are four bins of income06. As we can see above the predicted value decreases as income06 decreases.

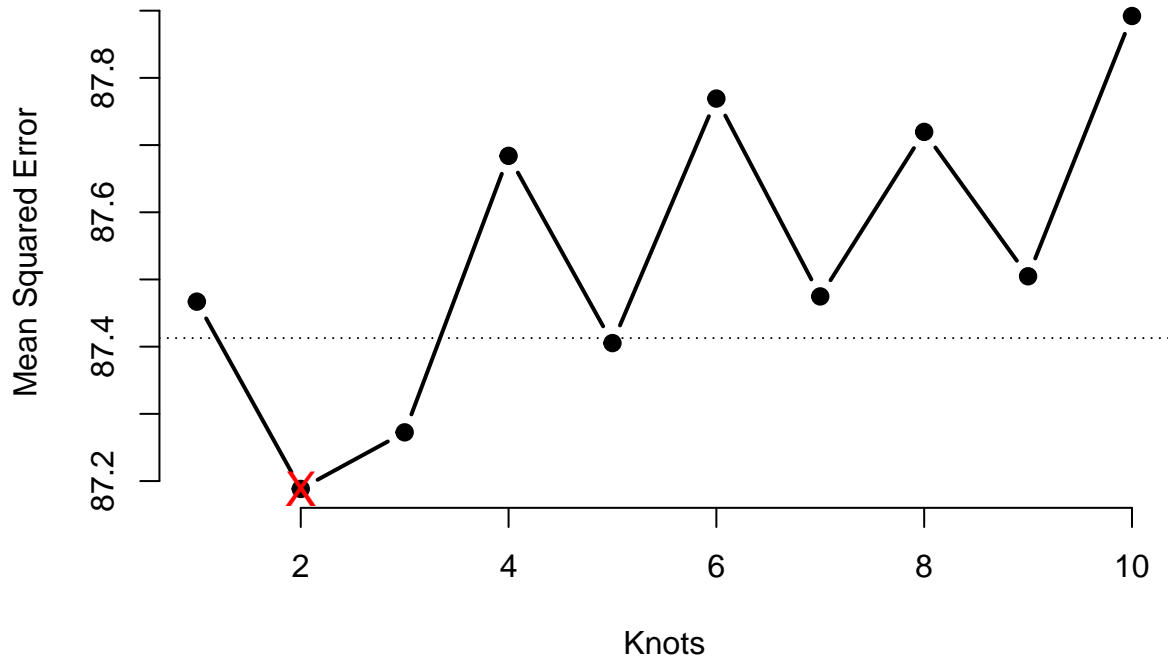
3. Fit a natural regression spline to predict egalit_scale as a function of income06. Use 10-fold cross-validation to select the optimal number of degrees of freedom, and present the results of the optimal model.

```
library(splines)
cv.MSE <- NA
for (i in 1:10) {
  glm.fit <- glm(egalit_scale ~ ns(income06, df = i), data = gss_train)
  cv.MSE[i] <- cv.glm(gss_train, glm.fit, K = 10)$delta[1]
}
table <- cbind(1:10, cv.MSE) %>%
  as.data.frame() %>%
  arrange(cv.MSE)
table
```

```
##      V1    cv.MSE
## 1     2 87.18831
## 2     3 87.27250
## 3     5 87.40533
## 4     1 87.46702
## 5     7 87.47478
## 6     9 87.50475
## 7     4 87.68397
## 8     8 87.71962
## 9     6 87.76919
## 10    10 87.89205
```

```
plot(1:10, cv.MSE, xlab = "Knots", ylab = "Mean Squared Error",
     type = "b", pch = 19, lwd = 2, bty = "n")
abline(h = min(cv.MSE, na.rm = TRUE) + sd(cv.MSE, na.rm = TRUE), lty = "dotted")

# highlight minimum
points(x = which.min(cv.MSE), y = min(cv.MSE, na.rm = TRUE), col = "red", pch = "X", cex = 1.5)
```



From above, the optimal number of knots is 2 knots.

```
best_natreg <- lm(egalit_scale ~ ns(income06, df = 2), data = gss_train)
plot_best_natreg <- cplot(best_natreg, "income06", draw=FALSE)
```

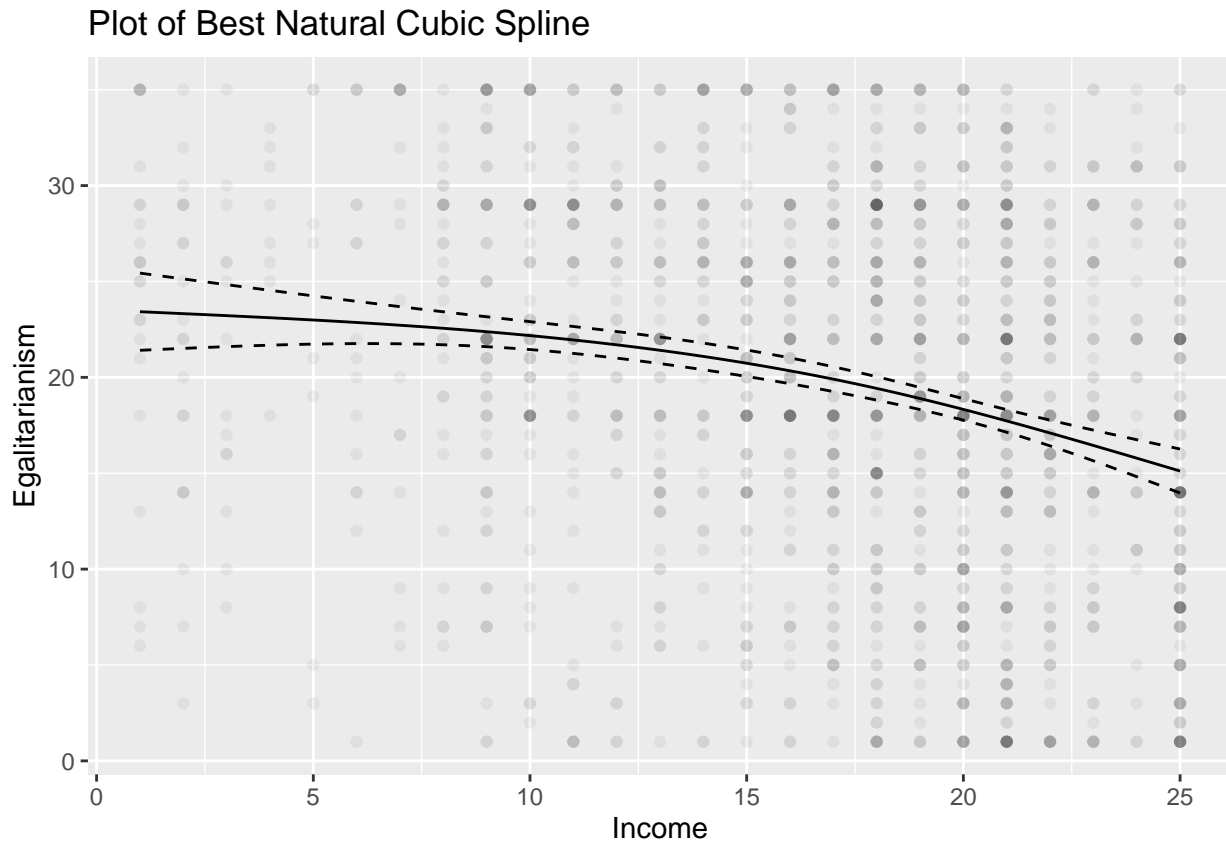
```
##      xvals      yvals      upper      lower
## 1      1 23.42271 25.43556 21.40987
## 2      2 23.32167 25.13244 21.51090
## 3      3 23.21783 24.83318 21.60248
## 4      4 23.10840 24.53808 21.67872
## 5      5 22.99058 24.24784 21.73332
## 6      6 22.86158 23.96357 21.75959
## 7      7 22.71860 23.68669 21.75050
## 8      8 22.55884 23.41854 21.69914
## 9      9 22.37951 23.15925 21.59977
## 10     10 22.17782 22.90625 21.44938
## 11     11 21.95096 22.65303 21.24888
## 12     12 21.69614 22.38954 21.00274
## 13     13 21.41057 22.10412 20.71702
## 14     14 21.09144 21.78564 20.39725
## 15     15 20.73597 21.42475 20.04720
## 16     16 20.34136 21.01450 19.66822
## 17     17 19.90481 20.55081 19.25881
## 18     18 19.42353 20.03365 18.81340
## 19     19 18.89631 19.47165 18.32097
## 20     20 18.32835 18.88975 17.76695
```

```
glm(egalit_scale ~ ns(income06, df = 2), data = gss_train) %>%
  cplot("income06", what = "prediction", n = 100, draw = FALSE) %>%
  ggplot(aes(x = xvals)) +
  geom_line(aes(y = yvals)) + # predicted value
  geom_line(aes(y = upper), linetype = 2) + # upper SE
  geom_line(aes(y = lower), linetype = 2) + # lower SE
  geom_point(data = gss_train, aes(income06, egalit_scale), alpha = 0.05) +
```



```
labs(x = "Income",
     y = "Egalitarianism",
     title = "Plot of Best Natural Cubic Spline")
```

##	xvals	yvals	upper	lower
## 1	1.000000	23.42271	25.43556	21.40987
## 2	1.242424	23.39832	25.36173	21.43492
## 3	1.484848	23.37390	25.28812	21.45967
## 4	1.727273	23.34939	25.21474	21.48404
## 5	1.969697	23.32476	25.14157	21.50794
## 6	2.212121	23.29997	25.06864	21.53131
## 7	2.454545	23.27499	24.99593	21.55404
## 8	2.696970	23.24976	24.92345	21.57607
## 9	2.939394	23.22426	24.85121	21.59731
## 10	3.181818	23.19843	24.77920	21.61766
## 11	3.424242	23.17225	24.70745	21.63705
## 12	3.666667	23.14567	24.63595	21.65539
## 13	3.909091	23.11865	24.56472	21.67258
## 14	4.151515	23.09116	24.49377	21.68854
## 15	4.393939	23.06314	24.42311	21.70317
## 16	4.636364	23.03457	24.35276	21.71638
## 17	4.878788	23.00540	24.28273	21.72807
## 18	5.121212	22.97560	24.21304	21.73815
## 19	5.363636	22.94511	24.14371	21.74652
## 20	5.606061	22.91391	24.07475	21.75307



From the above plot, we can see that there is a negative relationship between income06 and egalit_scale. The

higher an individual's income, the less egalitarian they are likely to be, which is similar to what we found out in the polynomial regression model.

4. Estimate the following models using all the available predictors (be sure to perform appropriate data pre-processing (e.g., feature standardization) and hyperparameter tuning (e.g. lambda for PCR/PLS, lambda and alpha for elastic net). Also use 10-fold cross-validation for each model to estimate the model's performance using MSE):

```
gss_train <- gss_train %>%  
  mutate_if(is.integer, scale) %>%  
  mutate_if(is.integer, c)  
gss_test <- gss_test %>%  
  mutate_if(is.integer, scale) %>%  
  mutate_if(is.integer, c)
```

- a. Linear regression

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
##
```

```
## Attaching package: 'Matrix'
```

```
## The following objects are masked from 'package:tidyr':
```

```
##
```

```
##      expand, pack, unpack
```

```
## Loaded glmnet 3.0-2
```

```
train.control <- trainControl(method = "cv", number = 10)  
lmodel <- train(egalit_scale ~ ., data = gss_train, method = "lm",  
               trControl = train.control)
```

```
lmodel
```

```
## Linear Regression
```

```
##
```

```
## 1481 samples
```

```
## 45 predictor
```

```
##
```

```
## No pre-processing
```

```
## Resampling: Cross-Validated (10 fold)
```

```
## Summary of sample sizes: 1333, 1334, 1334, 1333, 1334, 1332, ...
```

```
## Resampling results:
```

```
##
```

```
##      RMSE      Rsquared    MAE
```

```
## 0.8311914 0.3187364 0.6581222
```

```
##
```

```
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

```
MSE is 0.68196
```

- b. Elastic net regression

```
elr = train(  
  egalit_scale ~ ., data = gss_train,  
  method = "glmnet",  
  trControl = train.control,  
  tuneLength = 10
```

```

)
elr

## glmnet
##
## 1481 samples
## 45 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1333, 1332, 1332, 1332, 1334, 1334, ...
## Resampling results across tuning parameters:
##
##  alpha  lambda      RMSE      Rsquared  MAE
##  0.1    0.0002201401  0.8365024  0.3140697  0.6604380
##  0.1    0.0005085521  0.8364392  0.3141377  0.6603988
##  0.1    0.0011748214  0.8353316  0.3153555  0.6596767
##  0.1    0.0027139898  0.8336774  0.3171557  0.6584219
##  0.1    0.0062696687  0.8316225  0.3193209  0.6570465
##  0.1    0.0144837482  0.8284939  0.3225750  0.6548755
##  0.1    0.0334593377  0.8225068  0.3295103  0.6508953
##  0.1    0.0772954117  0.8160702  0.3379655  0.6480541
##  0.1    0.1785624307  0.8111845  0.3492419  0.6488689
##  0.1    0.4125023847  0.8229573  0.3487924  0.6668048
##  0.2    0.0002201401  0.8370554  0.3134894  0.6607432
##  0.2    0.0005085521  0.8360511  0.3145854  0.6601440
##  0.2    0.0011748214  0.8344648  0.3163433  0.6590347
##  0.2    0.0027139898  0.8324672  0.3185328  0.6575864
##  0.2    0.0062696687  0.8298971  0.3212793  0.6557937
##  0.2    0.0144837482  0.8246404  0.3273564  0.6520976
##  0.2    0.0334593377  0.8178312  0.3358551  0.6482994
##  0.2    0.0772954117  0.8105222  0.3478099  0.6456388
##  0.2    0.1785624307  0.8139078  0.3514000  0.6551488
##  0.2    0.4125023847  0.8480193  0.3215197  0.6942420
##  0.3    0.0002201401  0.8368357  0.3137582  0.6605736
##  0.3    0.0005085521  0.8357000  0.3149880  0.6599231
##  0.3    0.0011748214  0.8337531  0.3171558  0.6584854
##  0.3    0.0027139898  0.8315345  0.3195866  0.6569496
##  0.3    0.0062696687  0.8280705  0.3234515  0.6543909
##  0.3    0.0144837482  0.8216980  0.3311345  0.6501077
##  0.3    0.0334593377  0.8147713  0.3403667  0.6469221
##  0.3    0.0772954117  0.8089562  0.3522843  0.6463119
##  0.3    0.1785624307  0.8223913  0.3432373  0.6658919
##  0.3    0.4125023847  0.8719852  0.2917205  0.7188842
##  0.4    0.0002201401  0.8366788  0.3139102  0.6605352
##  0.4    0.0005085521  0.8352242  0.3155264  0.6595792
##  0.4    0.0011748214  0.8331702  0.3178083  0.6580294
##  0.4    0.0027139898  0.8308010  0.3204064  0.6564107
##  0.4    0.0062696687  0.8262456  0.3256919  0.6530565
##  0.4    0.0144837482  0.8194862  0.3340117  0.6488347
##  0.4    0.0334593377  0.8119668  0.3449166  0.6453445
##  0.4    0.0772954117  0.8101266  0.3524262  0.6491239
##  0.4    0.1785624307  0.8348690  0.3264407  0.6796028
##  0.4    0.4125023847  0.8923065  0.2641063  0.7386402

```

##	0.5	0.0002201401	0.8364301	0.3141883	0.6603726
##	0.5	0.0005085521	0.8347722	0.3160408	0.6592685
##	0.5	0.0011748214	0.8325633	0.3185001	0.6576304
##	0.5	0.0027139898	0.8300065	0.3213157	0.6558275
##	0.5	0.0062696687	0.8247269	0.3275586	0.6520282
##	0.5	0.0144837482	0.8178056	0.3362407	0.6480486
##	0.5	0.0334593377	0.8098879	0.3486664	0.6444581
##	0.5	0.0772954117	0.8122676	0.3511004	0.6526047
##	0.5	0.1785624307	0.8471493	0.3093296	0.6928375
##	0.5	0.4125023847	0.9080094	0.2445916	0.7542887
##	0.6	0.0002201401	0.8362234	0.3144231	0.6602430
##	0.6	0.0005085521	0.8344032	0.3164504	0.6589953
##	0.6	0.0011748214	0.8320783	0.3190503	0.6573308
##	0.6	0.0027139898	0.8292202	0.3222342	0.6552451
##	0.6	0.0062696687	0.8233625	0.3292783	0.6511052
##	0.6	0.0144837482	0.8163779	0.3382229	0.6474394
##	0.6	0.0334593377	0.8088245	0.3509641	0.6444379
##	0.6	0.0772954117	0.8158192	0.3472845	0.6574573
##	0.6	0.1785624307	0.8573568	0.2954940	0.7037562
##	0.6	0.4125023847	0.9188464	0.2416608	0.7654113
##	0.7	0.0002201401	0.8360678	0.3145983	0.6601475
##	0.7	0.0005085521	0.8340422	0.3168661	0.6587134
##	0.7	0.0011748214	0.8317129	0.3194590	0.6570665
##	0.7	0.0027139898	0.8283890	0.3232303	0.6545898
##	0.7	0.0062696687	0.8221309	0.3308409	0.6502999
##	0.7	0.0144837482	0.8150606	0.3401370	0.6468107
##	0.7	0.0334593377	0.8084297	0.3522845	0.6450185
##	0.7	0.0772954117	0.8206431	0.3411196	0.6631886
##	0.7	0.1785624307	0.8663703	0.2833476	0.7127077
##	0.7	0.4125023847	0.9310998	0.2356313	0.7771055
##	0.8	0.0002201401	0.8358555	0.3148370	0.6599982
##	0.8	0.0005085521	0.8337246	0.3172219	0.6584639
##	0.8	0.0011748214	0.8313738	0.3198390	0.6568084
##	0.8	0.0027139898	0.8275628	0.3242337	0.6539421
##	0.8	0.0062696687	0.8210328	0.3322407	0.6495911
##	0.8	0.0144837482	0.8138662	0.3419453	0.6461870
##	0.8	0.0334593377	0.8086136	0.3526818	0.6460694
##	0.8	0.0772954117	0.8261585	0.3333851	0.6693788
##	0.8	0.1785624307	0.8751481	0.2707595	0.7210004
##	0.8	0.4125023847	0.9426506	0.2336733	0.7873713
##	0.9	0.0002201401	0.8356514	0.3150631	0.6598637
##	0.9	0.0005085521	0.8334734	0.3175032	0.6582701
##	0.9	0.0011748214	0.8310528	0.3201980	0.6565807
##	0.9	0.0027139898	0.8267688	0.3252018	0.6533739
##	0.9	0.0062696687	0.8200743	0.3334672	0.6490388
##	0.9	0.0144837482	0.8125701	0.3439879	0.6454075
##	0.9	0.0334593377	0.8091234	0.3525833	0.6472806
##	0.9	0.0772954117	0.8319568	0.3248590	0.6757371
##	0.9	0.1785624307	0.8837063	0.2578570	0.7292716
##	0.9	0.4125023847	0.9556264	0.2336733	0.7982316
##	1.0	0.0002201401	0.8353801	0.3153765	0.6596813
##	1.0	0.0005085521	0.8331579	0.3178558	0.6580239
##	1.0	0.0011748214	0.8307369	0.3205500	0.6563472
##	1.0	0.0027139898	0.8260325	0.3261027	0.6528689

```
## 1.0 0.0062696687 0.8192339 0.3345506 0.6486226
## 1.0 0.0144837482 0.8113546 0.3459931 0.6446587
## 1.0 0.0334593377 0.8099367 0.3520357 0.6485929
## 1.0 0.0772954117 0.8375177 0.3167066 0.6818205
## 1.0 0.1785624307 0.8912788 0.2465169 0.7369840
## 1.0 0.4125023847 0.9708450 0.2336733 0.8123445
```

```
##
```

```
## RMSE was used to select the optimal model using the smallest value.
```

```
## The final values used for the model were alpha = 0.7 and lambda = 0.03345934.
```

The MSE with $\alpha = 0.6$ and $\lambda = 0.03345934$ is 0.64918, which is smaller than Linear Regression.

c. Principal component regression

```
pcr <- train(egalit_scale ~ ., data = gss_train, method = "pcr",
             trControl = train.control,
             metric = "RMSE", tuneLength = 50)
pcr
```

```
## Principal Component Analysis
```

```
##
```

```
## 1481 samples
```

```
## 45 predictor
```

```
##
```

```
## No pre-processing
```

```
## Resampling: Cross-Validated (10 fold)
```

```
## Summary of sample sizes: 1333, 1333, 1333, 1332, 1333, 1334, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

##	ncomp	RMSE	Rsquared	MAE
##	1	0.9996836	0.007171004	0.8379895
##	2	0.9367006	0.125604412	0.7743511
##	3	0.9211007	0.154603560	0.7619968
##	4	0.9172501	0.161305442	0.7592386
##	5	0.9160026	0.163426974	0.7572842
##	6	0.9166104	0.162234650	0.7575246
##	7	0.9134960	0.167888965	0.7538114
##	8	0.9120177	0.170717602	0.7509096
##	9	0.9126518	0.169666651	0.7517092
##	10	0.9127463	0.169880588	0.7509546
##	11	0.8981347	0.197925175	0.7387062
##	12	0.8989773	0.196823761	0.7397834
##	13	0.8726978	0.238594136	0.7058676
##	14	0.8716767	0.240528587	0.7043353
##	15	0.8658886	0.251184870	0.6969301
##	16	0.8650553	0.252962019	0.6963689
##	17	0.8652499	0.252707915	0.6960416
##	18	0.8621964	0.257247505	0.6942521
##	19	0.8597529	0.261482161	0.6935778
##	20	0.8591647	0.262432010	0.6918476
##	21	0.8477705	0.281484341	0.6797883
##	22	0.8483754	0.280433506	0.6805776
##	23	0.8445453	0.286892667	0.6767306
##	24	0.8454402	0.285692409	0.6775026
##	25	0.8457528	0.285471845	0.6767809
##	26	0.8438522	0.289073244	0.6742334

```
## 27 0.8413847 0.293930418 0.6719234
## 28 0.8359354 0.302673455 0.6683623
## 29 0.8358956 0.303329804 0.6672151
## 30 0.8369368 0.301957629 0.6669607
## 31 0.8246537 0.321995194 0.6560556
## 32 0.8254945 0.320856880 0.6564826
## 33 0.8247774 0.322011026 0.6558780
## 34 0.8250309 0.321632024 0.6564038
## 35 0.8237395 0.323671980 0.6555391
## 36 0.8226679 0.325530413 0.6542066
## 37 0.8216073 0.327180083 0.6531404
## 38 0.8234569 0.324406718 0.6535599
## 39 0.8230071 0.325259578 0.6536168
## 40 0.8217102 0.327348454 0.6533278
## 41 0.8208202 0.328681904 0.6528992
## 42 0.8199218 0.330330478 0.6523966
## 43 0.8172985 0.334506987 0.6503641
## 44 0.8175953 0.334386990 0.6501202
## 45 0.8179146 0.334227812 0.6510198
## 46 0.8187544 0.332924164 0.6514097
## 47 0.8195342 0.331887785 0.6517854
## 48 0.8198212 0.331426902 0.6517961
## 49 0.8198976 0.331329456 0.6519498
## 50 0.8192549 0.332423125 0.6514422
##
```

RMSE was used to select the optimal model using the smallest value.

The final value used for the model was ncomp = 43.

The best number of principal components is 45. The MSE is 0.66869 which is worse than Elastic Net but better than Linear Regression.

d. Partial least squares regression

```
pls = train(
  egalit_scale ~ ., data = gss_train,
  method = "pls",
  trControl = train.control,
  tuneLength = 50
)
pls
```

Partial Least Squares

##

1481 samples

45 predictor

##

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 1334, 1334, 1334, 1332, 1332, 1333, ...

Resampling results across tuning parameters:

##

##	ncomp	RMSE	Rsquared	MAE
##	1	0.8777560	0.2329988	0.7145092
##	2	0.8431473	0.2936222	0.6722379
##	3	0.8315987	0.3122365	0.6624590
##	4	0.8246272	0.3233038	0.6524952

```

##      5      0.8220281  0.3285686  0.6523111
##      6      0.8246923  0.3257007  0.6537868
##      7      0.8281283  0.3214520  0.6557273
##      8      0.8308319  0.3185297  0.6572824
##      9      0.8317398  0.3177886  0.6569875
##     10      0.8323539  0.3175554  0.6572395
##     11      0.8321146  0.3180799  0.6571540
##     12      0.8301190  0.3213693  0.6545401
##     13      0.8292263  0.3229184  0.6537179
##     14      0.8283409  0.3241549  0.6525641
##     15      0.8285332  0.3239940  0.6528852
##     16      0.8288595  0.3236448  0.6529116
##     17      0.8293823  0.3231619  0.6531783
##     18      0.8297630  0.3226391  0.6537405
##     19      0.8304124  0.3217059  0.6541253
##     20      0.8304573  0.3217161  0.6541223
##     21      0.8309551  0.3211280  0.6545872
##     22      0.8307994  0.3213349  0.6545620
##     23      0.8313209  0.3207078  0.6553449
##     24      0.8312885  0.3207917  0.6549490
##     25      0.8313487  0.3207471  0.6546072
##     26      0.8316463  0.3203758  0.6550663
##     27      0.8317374  0.3202606  0.6549979
##     28      0.8317954  0.3201621  0.6550590
##     29      0.8317867  0.3201634  0.6550650
##     30      0.8320297  0.3198300  0.6552192
##     31      0.8320539  0.3197894  0.6552900
##     32      0.8319783  0.3198792  0.6553162
##     33      0.8317907  0.3201312  0.6551271
##     34      0.8317480  0.3201731  0.6550332
##     35      0.8317695  0.3201573  0.6550276
##     36      0.8317809  0.3201406  0.6549899
##     37      0.8319677  0.3199133  0.6551349
##     38      0.8320562  0.3198192  0.6552136
##     39      0.8323638  0.3193987  0.6554692
##     40      0.8326703  0.3189718  0.6557165
##     41      0.8328532  0.3187369  0.6557931
##     42      0.8331943  0.3182776  0.6559403
##     43      0.8331117  0.3183879  0.6559335
##     44      0.8329476  0.3186932  0.6558402
##     45      0.8324048  0.3194773  0.6554411
##     46      0.8324381  0.3194680  0.6556794
##     47      0.8325683  0.3193471  0.6555748
##     48      0.8329309  0.3188085  0.6557993
##     49      0.8331389  0.3185127  0.6559339
##     50      0.8332875  0.3182699  0.6560132
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was ncomp = 5.

```

The best number of principal components is 5 which is much smaller than Principal Component Analysis. The MSE is 0.66918 which is larger than Principal Component Analysis but still better than Linear Regression.

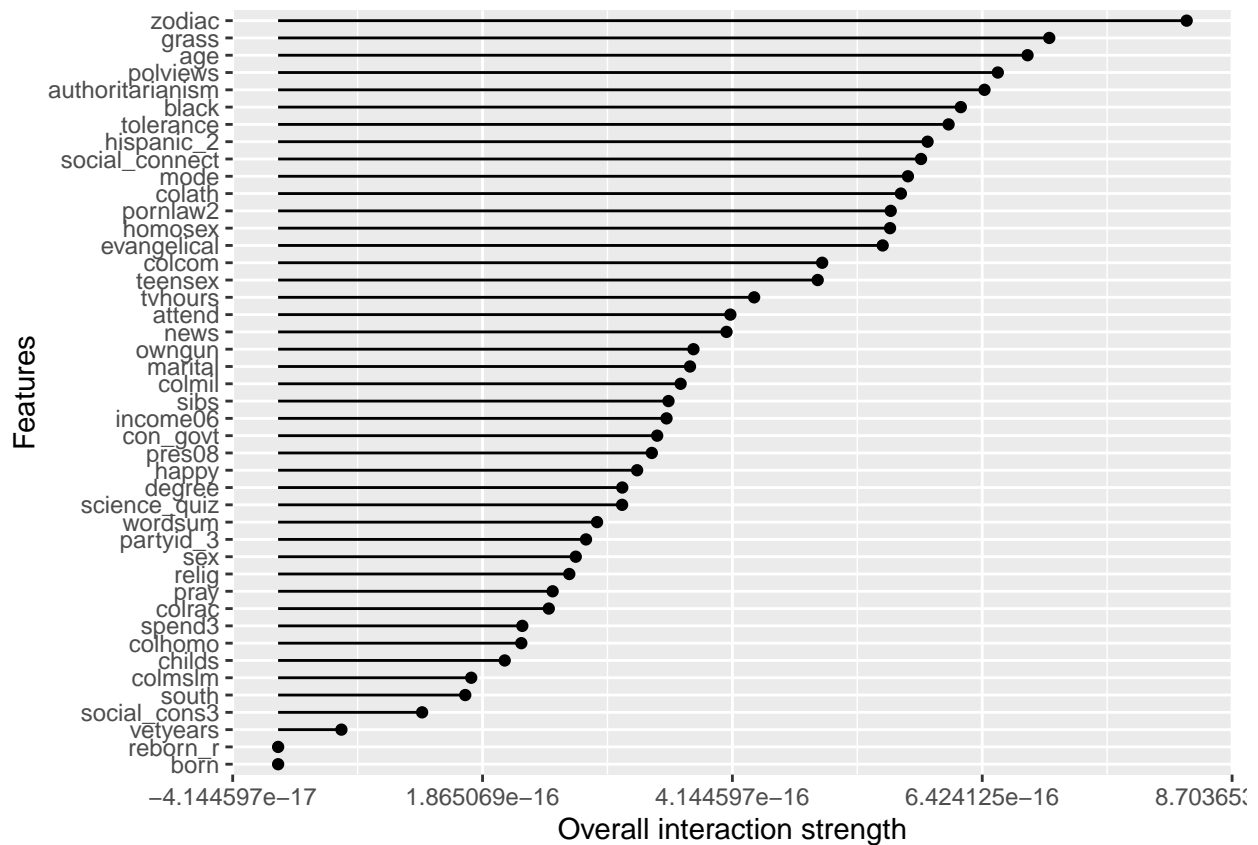
Elastic Net is the best model for this dataset when comparing the MSE.

5. For each final tuned version of each model fit, evaluate feature importance by generating feature interaction plots. Upon visual presentation, be sure to discuss the substantive results for these models and in comparison to each other (e.g., talk about feature importance, conditional effects, how these are ranked differently across different models, etc.)

```
gss_train <- read.csv('~/Desktop/problem-set-4-master/data/gss_train.csv')
gss_test <- read.csv('~/Desktop/problem-set-4-master/data/gss_test.csv')
lmodel <- train(egalit_scale ~., data = gss_train, method = "lm",
               trControl = train.control)
elr = train(egalit_scale ~ ., data = gss_train,
            method = "glmnet",
            trControl = train.control,
            tuneLength = 10)
pcr <- train(egalit_scale ~ ., data = gss_train, method = "pcr",
            trControl = train.control,
            metric = "RMSE", tuneLength = 50)
pls = train(egalit_scale ~ ., data = gss_train,
            method = "pls",
            trControl = train.control,
            tuneLength = 50)
```


```
library(iml)
features <- gss_test %>%
  dplyr::select(- egalit_scale)
response <- as.integer(gss_test$egalit_scale)
```

```
pred_lm <- Predictor$new(
  model = lmodel,
  data = features,
  y = response
)
interact_lm <- Interaction$new(pred_lm)
plot(interact_lm)
```


In this linear model, polviews, homosexuality, and income06 are the three features with the highest interaction strength.

```
pred_elr <- Predictor$new(
  model = elr,
  data = features,
  y = response
)
interact_elr <- Interaction$new(pred_elr)
plot(interact_elr)
```




In this elastic net model, relig, income06, and vetyears are the three features with the highest interaction strength.

```
pred_pcr <- Predictor$new(
  model = pcr,
  data = features,
  y = response
)
interact_pcr <- Interaction$new(pred_pcr)
plot(interact_pcr)
```



In this PCR model, attend, social_cons3, and marital are the three features with the highest interaction strength.

```
pred_pls <- Predictor$new(
  model = pls,
  data = features,
  y = response
)
interact_pls <- Interaction$new(pred_pls)
plot(interact_pls)
```



In this PLS model, evangelical, colmslm and pres08 are the three features with the highest interaction strength.

All models have different rank of important features by interaction strength because they prioritize and penalize different things.