

## Homework 4: Moving Beyond Linearity

Tianyue Niu

```
In [517]: #import necessary packages
import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import cross_val_score
from patsy import dmatrix
import matplotlib.pyplot as plt
import statsmodels.api as sm
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import ElasticNet
from sklearn.decomposition import PCA
from sklearn.cross_decomposition import PLSRegression
from sklearn.inspection import plot_partial_dependence
from mlxtend.evaluate import feature_importance_permutation
```

```
In [518]: #read in data
test = pd.read_csv('gss_test.csv')
train = pd.read_csv('gss_train.csv')
```

### Egalitarianism and income

1)(20 points) Perform polynomial regression to predict `egalit_scale` as a function of `income06`. Use and plot 10-fold cross-validation to select the optimal degree  $d$  for the polynomial based on the MSE. Plot the resulting polynomial fit to the data, and also graph the average marginal effect (AME) of `income06` across its potential values. Be sure to provide substantive interpretation of the results.

```
In [519]: #Didn't use sklearn PolynomialFeatures because the function automatically includes inter
          action
          #terms in the regression fit and these interaction terms can't be removed.
```

```
In [520]: #define X and y
X = np.array(train['income06']).reshape(-1,1)
y = np.array(train['egalit_scale']).reshape(-1,1)
```

```

In [521]: def find_best_deg(X, y):
            cv = {}
            min_error = 10000000
            min_deg = 0
            degrees = np.arange(1, 21)
            x = pd.DataFrame()

            for deg in degrees:
                #manually add polynomials of X to x
                x[deg] = X ** deg
                # Fit polynomial regression
                lr = LinearRegression()
                #get error
                error = np.sum(-cross_val_score(lr, x, y, cv=10, scoring='neg_mean_squared_error'))/10
                cv[deg] = error
                if error < min_error:
                    min_error = error
                    min_deg = deg

            return min_deg, cv

```

```

In [522]: find_best_deg(train['income06'], y)


```

```

Out[522]: (12,
           {1: 87.96300193115194,
            2: 87.65085217286715,
            3: 87.7154238094978,
            4: 87.7506412125094,
            5: 87.97322944676486,
            6: 88.01923052537474,
            7: 87.70981059410886,
            8: 87.73114625218366,
            9: 87.61920345310567,
            10: 87.51660188653618,
            11: 87.5770963856581,
            12: 87.42373138398571,
            13: 87.44309717782467,
            14: 87.53385243089564,
            15: 87.58130024172851,
            16: 87.64923123298351,
            17: 87.69882301886456,
            18: 87.70242260890136,
            19: 87.75905670902408,
            20: 87.84239432606995})

```

```
In [523]: #Plot the MSE graph
lists = sorted(find_best_deg(train['income06'],y)[1].items())
degree, mse = zip(*lists)
plt.figure(figsize=(10,5))
plt.plot(degree, mse)
plt.xlabel('Degree of polynomials')
plt.ylabel('Cross validated MSE')
plt.title('Model MSE vs model degree of polynomials');
```



We see that the MSE fluctuates as degree of polynomials increases. The cross-validated MSE is lowest at degree = 12. There also exists a local minimum at degree = 2.

```
In [524]: def calculate_ame (X, y, degree):

    x = pd.DataFrame()
    x_drv = pd.DataFrame()
    ame_total = []

    for deg in range(degree+1):
        if deg != 0:
            x[deg] = X ** deg
        lm = LinearRegression().fit(x, y)
        coefs = lm.coef_

        for deg in range(degree):
            x_drv[deg] = (deg+1)* coefs[deg] * (X ** deg)
            ame_x = x_drv[deg].mean()
            ame_total.append(ame_x)
        ame = np.sum(ame_total)

    return ame, coefs
```

```
In [525]: ame, coef_degree_12 = calculate_ame(train['income06'], train['egalit_scale'], 12)
```

```
In [526]: coef_degree_12
```

```
Out[526]: array([ 8.33052123e-06,  3.86413357e-05,  2.27267669e-04,  7.25124016e-04,
 1.22065917e-03, -6.62426809e-04,  1.34765206e-04, -1.45754642e-05,
 9.20230578e-07, -3.41025761e-08,  6.88602186e-10, -5.85313884e-12])
```

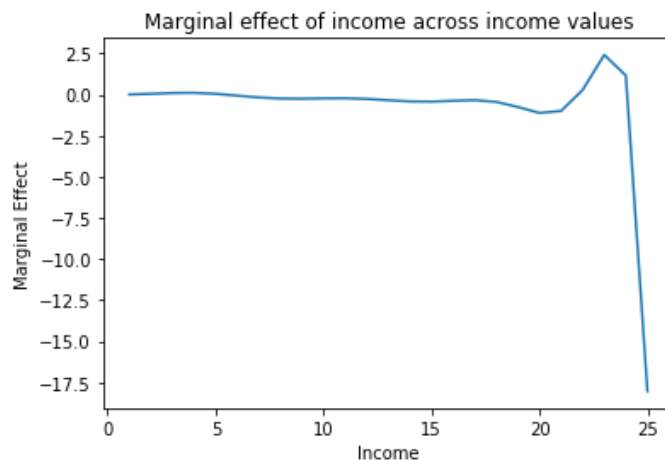
```
In [527]: print('The AME of income is {}'.format(ame))
```

The AME of income is -1.837827013750939

```
In [528]: #graph marginal effect
def get_marginal_effect(x):
    y = 0
    for i in range(12):
        y += coef_degree_12[i] * (i+1) * (x ** (i))
    return y

marginal_effects = []
for x in range(1,27):
    marginal_effects.append(get_marginal_effect(x))
```

```
In [529]: income_values = np.arange(1,26)
marginal_income = pd.DataFrame(zip(list(income_values), marginal_effects),
                               columns = ['income_value', 'marginal_effect'])
plt.plot(marginal_income['income_value'], marginal_income['marginal_effect'])
plt.title('Marginal effect of income across income values')
plt.xlabel('Income')
plt.ylabel('Marginal Effect');
```



2)(20 points) Fit a step function to predict `egalit_scale` as a function of `income06`, and perform 10-fold cross-validation to choose the optimal number of cuts. Plot the fit and interpret the results.

```
In [530]: def find_best_bin(bins = np.arange(2, 11)):

    cv = {}
    min_error = 10000000
    best_bin = 0

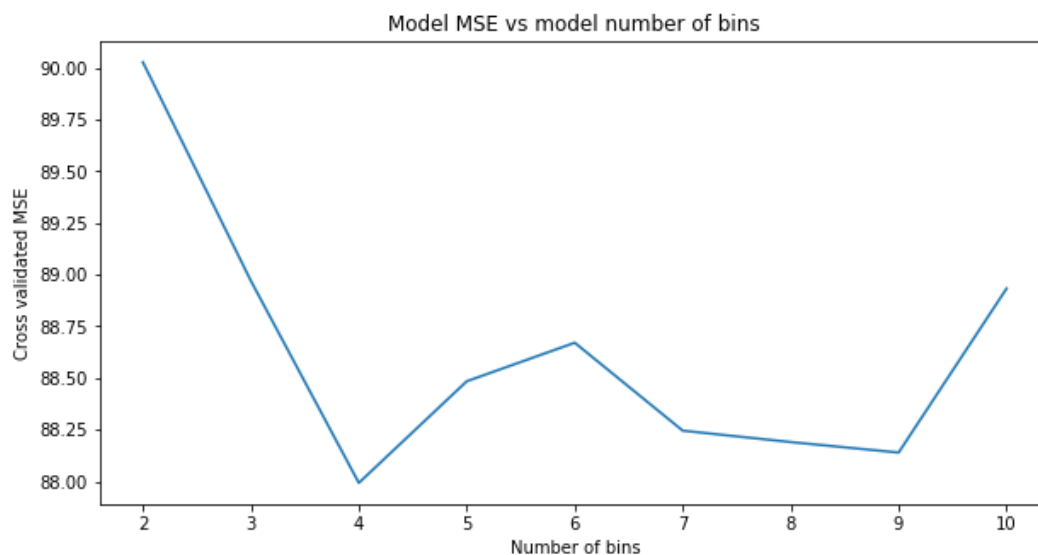
    for bin in bins:
        #cut data into bins
        step, bins = pd.cut(train.income06, bins=bin, retbins = True)
        #create dummy variables
        steps_dummies = pd.get_dummies(step)
        #combine y and x into one data frame
        step_df = pd.concat([train.egalit_scale, steps_dummies], axis = 1)
        x = np.array(step_df[step_df.columns[2:]])
        y = np.array(step_df['egalit_scale'])
        lr = LinearRegression()
        error = np.sum(-cross_val_score(lr, x , y, cv=10, scoring='neg_mean_squared_erro
r'))/10
        cv[bin] = error
        if error < min_error:
            min_error = error
            best_bin = bin

    return best_bin, cv
```

```
In [531]: find_best_bin()
```

```
Out[531]: (4,
           {2: 90.02781799497726,
            3: 88.97028485125078,
            4: 87.99268202181436,
            5: 88.48448329543504,
            6: 88.67079198404846,
            7: 88.24549422096784,
            8: 88.19032257110334,
            9: 88.13960580789443,
            10: 88.93320277814779})
```

```
In [532]: #Plot the MSE graph
bins_list = sorted(find_best_bin()[1].items())
bins, mse = zip(*bins_list)
plt.figure(figsize=(10,5))
plt.plot(bins, mse)
plt.xlabel('Number of bins')
plt.ylabel('Cross validated MSE')
plt.title('Model MSE vs model number of bins');
```



The above graph shows that the optimal number of bins = 4. This means that when we break income into four different groups, the linear regression model will predict `egalit_scale` with the lowest error.

3)(20 points) Fit a natural regression spline to predict `egalit_scale` as a function of `income06`. Use 10-fold cross-validation to select the optimal number of degrees of freedom, and present the results of the optimal model.

```

In [533]: def find_best_deg(X, y):
            cv = {}
            min_error = 10000000
            min_deg = 0
            degrees = np.arange(3, 21)
            x = 0

            for deg in degrees:
                #transform to cubic spline
                x = dmatrix("cr(X,df = {})".format(deg), {"X": X}, return_type='dataframe')
                # Fit regression
                lr = LinearRegression()
                #get error
                error = np.sum(-cross_val_score(lr, x, y, cv=10, scoring='neg_mean_squared_error'))/10
                cv[deg] = error
                if error < min_error:
                    min_error = error
                    min_deg = deg

            return min_deg, cv

```

```

In [534]: find_best_deg(X,y)


```

```

Out[534]: (3,
           {3: 87.66117598417557,
            4: 87.7027438611445,
            5: 87.70523439574966,
            6: 88.10043282440009,
            7: 88.09833195179264,
            8: 88.20904399518508,
            9: 88.26253460814083,
            10: 88.31086968848001,
            11: 88.16429451853226,
            12: 88.11220442409994,
            13: 88.15088832132531,
            14: 88.43623469335596,
            15: 88.42653388451811,
            16: 88.39292563579053,
            17: 88.45039559291578,
            18: 88.48043046470737,
            19: 88.5637800942234,
            20: 88.44447802979225})

```

```
In [535]: #Plot the MSE graph
spline_list = sorted(find_best_deg(X,y)[1].items())
degree, mse = zip(*spline_list)
plt.figure(figsize=(10,5))
plt.plot(degree, mse)
plt.xlabel('Degrees of freedom')
plt.ylabel('Cross validated MSE')
plt.title('Model MSE vs models degree of freedom');
```



The optimal model is when degree of freedom = 3. The MSE of that model is 87.66117598417557.

### Egalitarianism and everything

4)(20 points total) Estimate the following models using all the available predictors (be sure to perform appropriate data pre-processing (e.g., feature standardization) and hyperparameter tuning (e.g. lambda for PCR/PLS, lambda and alpha for elastic net). Also use 10-fold cross-validation for each model to estimate the model's performance using MSE):

```
In [536]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1481 entries, 0 to 1480
Data columns (total 45 columns):
age                1481 non-null int64
attend             1481 non-null object
authoritarianism   1481 non-null int64
black              1481 non-null object
born               1481 non-null object
childs            1481 non-null int64
colath             1481 non-null object
colrac             1481 non-null object
colcom            1481 non-null object
colmil            1481 non-null object
colhomo           1481 non-null object
colmslm           1481 non-null object
con_govt          1481 non-null int64
degree            1481 non-null object
egalit_scale      1481 non-null int64
evangelical       1481 non-null object
grass             1481 non-null object
happy             1481 non-null object
hispanic_2        1481 non-null object
homosex           1481 non-null object
income06          1481 non-null int64
marital           1481 non-null object
mode              1481 non-null object
news              1481 non-null object
owngun            1481 non-null object
partyid_3         1481 non-null object
polviews          1481 non-null object
pornlaw2          1481 non-null object
pray              1481 non-null object
pres08            1481 non-null object
reborn_r          1481 non-null object
relig             1481 non-null object
science_quiz      1481 non-null int64
sex               1481 non-null object
sibs              1481 non-null int64
social_connect     1481 non-null int64
social_cons3      1481 non-null object
south             1481 non-null object
spend3            1481 non-null object
teensex           1481 non-null object
tolerance         1481 non-null int64
tvhours           1481 non-null int64
vetyears          1481 non-null object
wordsum           1481 non-null int64
zodiac            1481 non-null object
dtypes: int64(12), object(33)
memory usage: 520.7+ KB
```

```
In [537]: #Create dummies for categorical variables
```

```
def create_dummies (df, new_df):
    for predictor in df:
        if df[predictor].dtypes == object:
            temp = pd.get_dummies(df[predictor])
            temp = temp.drop(temp.columns[0], axis=1)
            new_df = pd.concat([new_df, temp], axis=1)
    return new_df
```



In [538]: *#Standardize data*

```
scaler = MinMaxScaler(feature_range=(0,1))

def standardize (df):
    new_df = pd.DataFrame()
    for predictor in df:
        if df[predictor].dtypes == 'int64':
            column = df[predictor].values.reshape(-1,1)
            scaler.fit(column)
            new_df[predictor] = scaler.transform(column).reshape(1,-1)[0]
    return new_df
```

In [539]: *#transform and standardize X and Y*

```
X_train = standardize(train)
y_train = X_train['egalit_scale']
X_train = create_dummies(train, X_train)
X_train.drop('egalit_scale', axis=1, inplace=True)
X_test = standardize(test)
y_test = X_test['egalit_scale']
X_test = create_dummies(test, X_test)
X_test.drop('egalit_scale', axis=1, inplace=True)
```

In [540]: X\_train.head(2)

Out[540]:

	age	authoritarianism	childs	con_govt	income06	science_quiz	sibs	social_connect	tolerance	tvhours
0	0.042254	0.571429	0.00	1.000000	1.000000	0.7	0.066667	0.416667	0.666667	0.125
1	0.338028	0.571429	0.25	0.333333	0.916667	1.0	0.033333	0.416667	0.866667	0.125

2 rows × 102 columns

a. (5 points) *Linear regression*

In [541]:

```
lr = LinearRegression()
np.sum(-cross_val_score(lr, X_train, y_train, cv=10, scoring='neg_mean_squared_error'))/
10
```

Out[541]: 0.05520732886554378

In [542]: lr = lr.fit(X\_train, y\_train)

b. (5 points) *Elastic net regression*

```
In [543]: #define a function to find the best lambda and alpha
def find_best_para():
    min_mse = 10000
    best_model = None
    best_lambda = None
    best_alpha = None

    for i in np.arange(0.1,1,0.1):
        for j in np.arange(0.1,1,0.1):
            elastic = ElasticNet(l1_ratio=i, alpha=j)
            mse = np.sum(-cross_val_score(elastic, X_train, y_train, cv=10,
                                          scoring='neg_mean_squared_error'))/10

            if mse < min_mse:
                min_mse = mse
                best_model = elastic
                best_lambda = i
                best_alpha = j

    return best_lambda, best_alpha, min_mse
```

```
In [544]: a,b,c = find_best_para()
```

```
In [545]: "The best model has lambda = {}, alpha={}, \
and the MSE of this model is equal to {}".format(a,b,c)
```

```
Out[545]: 'The best model has lambda = 0.1, alpha =0.1, and the MSE of this model is equal to 0.
05980777651982657'
```

```
In [546]: elastic = ElasticNet(l1_ratio=0.1, alpha=0.1).fit(X_train, y_train)
```

c. (5 points) *Principal component regression*

```
In [547]: def find_component_pca():
    min_mse = 10000
    number = 0
    for i in np.arange(1,21,1):
        pca = PCA(n_components = i)
        pc = pca.fit_transform(X_train)
        mse = np.sum(-cross_val_score(lr, pc, y_train, cv=10,
                                      scoring='neg_mean_squared_error'))/10

        if mse < min_mse:
            min_mse = mse
            number = i
    return i, min_mse
```

```
In [548]: find_component_pca()
```

```
Out[548]: (20, 0.05637811069577927)
```

```
In [549]: pca = PCA(n_components = 20)
pc = pca.fit_transform(X_train)
lr_pca = LinearRegression().fit(pc, y_train)
```

We see that the best model has components = 20 (choosing among 1-20) and MSE =0.056623.

d. (5 points) *Partial least squares regression*

```
In [550]: def find_component_pls():
min_mse = 10000
number = 0
for i in np.arange(1,21,1):
    pls = PLSRegression(n_components=i)
    mse = np.sum(-cross_val_score(pls, X_train, y_train, cv=10,
                                scoring='neg_mean_squared_error'))/10

    if mse < min_mse:
        min_mse = mse
        number = i
return i, min_mse
```

```
In [551]: find_component_pls()
```

```
Out[551]: (20, 0.05486714606062361)
```

We see that the best PLS model also has components = 20 (choosing among 1-20), its MSE =0.054867.

```
In [552]: pls = PLSRegression(n_components=20)
pls.fit(X_train, y_train)
```

```
Out[552]: PLSRegression(copy=True, max_iter=500, n_components=20, scale=True, tol=1e-06)
```

5)(20 points) For each final tuned version of each model fit, evaluate feature importance by generating feature interaction plots. Upon visual presentation, be sure to discuss the substantive results for these models and in comparison to each other (e.g., talk about feature importance, conditional effects, how these are ranked differently across different models, etc.).

```
In [624]: X_train.head(1)
```

```
Out[624]:
```


	age	authoritarianism	childs	con_govt	income06	science_quiz	sibs	social_connect	tolerance	tvhours
0	0.042254	0.571429	0.0	1.0	1.0	0.7	0.066667	0.416667	0.666667	0.125

1 rows × 102 columns


```
In [590]: features = []
for i in range(102):
    features.append((4,i))
```

```
In [630]: #There is no package for plotting feature interaction plot in python
#I will manually plot the interaction between
#income06 and 4 other variables for different models
```


```
In [628]: #linear regression
#0=age, 1=authoritarianism, 2= childs, 3=con_govt, 4 = income06
plot_partial_dependence(lr, X_train, [(4,0), (4,1)])
plot_partial_dependence(lr, X_train, [(4,2), (4,3)])
```



```
In [632]: #elastic net
#0=age, 1=authoritarianism, 2= childs, 3=con_govt, 4 = income06
plot_partial_dependence(elastic, X_train, [(4,0), (4,1)])
plot_partial_dependence(elastic, X_train, [(4,2), (4,3)])
```



```
In [633]: #pls
#0=age, 1=authoritarianism, 2= childs, 3=con_govt, 4 = income06
plot_partial_dependence(pls, X_train, [(4,0), (4,1)])
plot_partial_dependence(pls, X_train, [(4,2), (4,3)])
```



0=age, 1=authoritarianism, 2= childs, 3=con\_govt, 4 = income06.

We see that for pls and linear regression, there seems to exist interaction between income and age, income and childs, and income and congovt. No significant interaction is plotted for income and authoritarianism.

However, when using elastic net, we see that all interactions disappeared.