

# Homework 4: Moving Beyond Linearity

Wen Li Teng

February 16 2020

```
library(tidyverse)
library(caret)
library(margins)
library(splines)
library(ggplot2)
library(ggfortify)
library(robustHD)
library(glmnet)
library(pls)
library(iml)
library(h2o)

train <- read.csv("data/gss_train.csv")
test <- read.csv("data/gss_test.csv")

set.seed(1234)
```

## 1. POLYNOMIAL REGRESSION

Perform polynomial regression to predict `egalit_scale` as a function of `income06`

Use and plot 10-fold cross-validation

```
polynomial <- 1:10
poly_rsquared <- rep(0,10)
poly_RMSE <- rep(0,10)

train_control <- trainControl(method = "CV", number = 10)

for (i in 1:10) {
  poly_formula <- bquote(egalit_scale ~ poly(income06, .(i)))
  poly_mod <- train(as.formula(poly_formula),
                    data = train,
                    method = "lm",
                    trControl = train_control)
  poly_rsquared[i] <- poly_mod$results$Rsquared
  poly_RMSE[i] <- poly_mod$results$RMSE
}
```

```
poly_table <- cbind(polynomial, poly_rsquared, poly_RMSE) %>%
  as.data.frame() %>%
  arrange(poly_RMSE)
poly_table
```

```
##      polynomial poly_rsquared poly_RMSE
## 1             9    0.06625946  9.326062
## 2             2    0.06323179  9.327615
## 3             3    0.06172276  9.330483
## 4             8    0.06119367  9.332842
## 5             6    0.05792979  9.340852
## 6             5    0.06439519  9.340872
## 7             7    0.06328038  9.341814
## 8             4    0.06304707  9.343810
## 9            10    0.06131416  9.349980
## 10            1    0.06256509  9.350900
```

## Select optimal degree for the polynomial based on MSE

Following the intuition of James et al. (2013), choose a polynomial that balances test error (accuracy) with simplicity (interpretability). The 9th-order polynomial is the most accurate, but is the 2nd-least interpretable. The 2nd-order polynomial is the 2nd-most accurate and is the 2nd-most interpretable. Therefore, choose the 2nd-order polynomial.


## Plot polynomial fit to the data

```
poly_mod <- lm(egalit_scale ~ poly(income06, 2), data = train)
poly_pred <- predict(poly_mod, newdata = test)
poly_df <- data.frame(income = test$income06,
                      actual = test$egalit_scale,
                      predicted = poly_pred)

ggplot(poly_df, aes(x = income)) +
  geom_point(aes(y = actual, color = "actual")) +
  geom_point(aes(y = predicted, color = "predicted")) +
  geom_smooth(method = "lm",
              formula = y ~ poly(x, 2),
              aes(y = actual, color = "actual")) +
  geom_smooth(data = poly_df,
              method = "lm",
              formula = y ~ poly(x, 2),
              aes(x = income, y = predicted,
                  color = "predicted")) +
  scale_colour_manual("", values = c("actual" = "blue",
                                     "predicted" = "red")) +
  xlab("income06") + ylab("egalit_scale") +
  labs(title = "Predicted egalit_scale",
       subtitle = "Polynomial regression with 2nd-order polynomial")
```

## Predicted egalit\_scale

Polynomial regression with 2nd-order polynomial




Graph the average marginal effect (AME) of income06 across potential values

```
poly_dat <- cplot(poly_mod, "income06", draw = FALSE)
```

##	xvals	yvals	upper	lower
## 1	1	23.11631	25.28371	20.94890
## 2	2	23.15180	25.04001	21.26359
## 3	3	23.15524	24.79232	21.51817
## 4	4	23.12665	24.54195	21.71134
## 5	5	23.06600	24.29036	21.84165
## 6	6	22.97331	24.03899	21.90764
## 7	7	22.84858	23.78870	21.90846
## 8	8	22.69180	23.53894	21.84467
## 9	9	22.50298	23.28678	21.71917
## 10	10	22.28211	23.02670	21.53752
## 11	11	22.02920	22.75132	21.30708
## 12	12	21.74424	22.45297	21.03552
## 13	13	21.42724	22.12502	20.72946
## 14	14	21.07819	21.76262	20.39376
## 15	15	20.69710	21.36290	20.03130
## 16	16	20.28396	20.92501	19.64291
## 17	17	19.83878	20.45044	19.22712
## 18	18	19.36155	19.94356	18.77955

```
## 19    19 18.85228 19.41247 18.29210
## 20    20 18.31096 18.86919 17.75274
```

```
ggplot(poly_dat, aes(x = xvals)) +
  geom_line(aes(y = yvals)) +
  geom_line(aes(y = upper), linetype = 2) +
  geom_line(aes(y = lower), linetype = 2) +
  geom_hline(yintercept = 0) +
  ggtitle("AME of income06 on egalit_scale") +
  xlab("egalit_scale") + ylab("AME of income06")
```



## Interpret the results

The polynomial regression model (with a 2nd-order polynomial) is of middling accuracy, judging from the R-squared values of the model and the fit of the model to the data. Nonetheless, the model can still be informative. In particular, the model reflects an inverse relationship between the `income06` and `egalit_scale`. The average marginal effects of `income06` decreases as `egalit_scale` increases. Put differently, income appears to have less of an effect on egalitarianism at higher levels of income.

## 2. STEP FUNCTION

Fit a step function to predict `egalit_scale` as a function of `income06`

Perform 10-fold cross-validation to choose optimal number of cuts

```
cut <- 1:10
step_rsquared <- rep(0,10)
step_RMSE <- rep(0,10)

for (i in 1:10) {
  step_formula <- bquote(egalit_scale ~ cut(income06, .(i+1)))
  step_mod <- train(as.formula(step_formula),
                    data = train,
                    method = "lm",
                    trControl = train_control)
  step_rsquared[i] <- step_mod$results$Rsquared
  step_RMSE[i] <- step_mod$results$RMSE
}

step_table <- cbind(cut, step_rsquared, step_RMSE) %>%
  as.data.frame() %>%
  arrange(step_RMSE)
step_table
```

##	cut	step_rsquared	step_RMSE
## 1	3	0.06475192	9.340726
## 2	6	0.05863575	9.352234
## 3	7	0.05603677	9.365507
## 4	4	0.05684573	9.369439
## 5	8	0.06325533	9.373954
## 6	5	0.05814558	9.391172
## 7	10	0.05184705	9.393802
## 8	9	0.05010242	9.402859
## 9	2	0.04826193	9.407046
## 10	1	0.04640243	9.471710

Select 3 as number of cuts, given that it has the lowest RMSE (accuracy) and is a relatively low number of cuts (interpretability).

Plot the fit

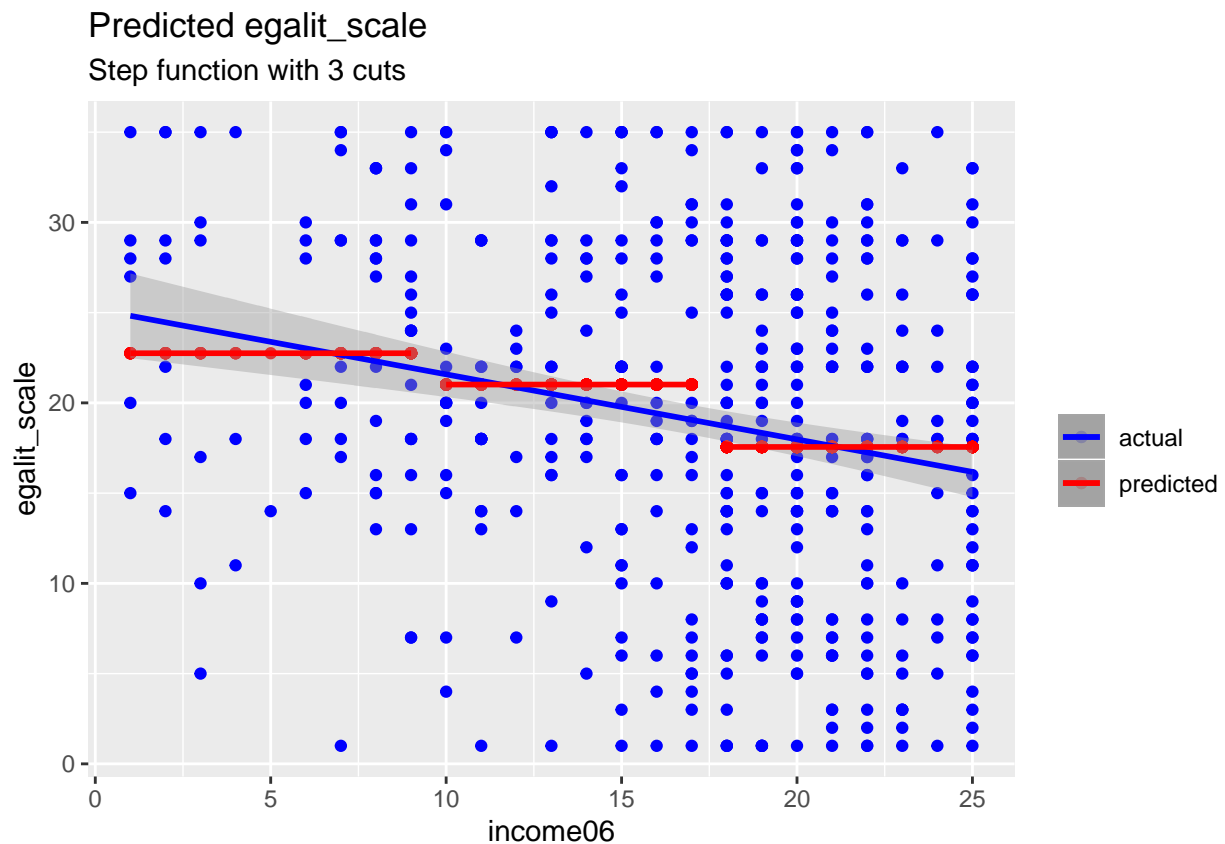
```
step_mod <- lm(egalit_scale ~ cut(income06, 3), data = train)
step_pred <- predict(step_mod, newdata = test)
step_df <- data.frame(income = test$income06,
                      actual = test$egalit_scale,
                      predicted = step_pred)
step_df1 <- step_df %>%
  filter(income <= 9)
```

```

step_df2 <- step_df %>%
  filter(income > 9 & income <= 17)
step_df3 <- step_df %>%
  filter(income > 17)

ggplot(step_df, aes (x = income)) +
  geom_point(aes(y = actual, color = "actual")) +
  geom_point(aes(y = predicted, color = "predicted")) +
  geom_smooth(method = "lm", aes(y = actual, color = "actual")) +
  geom_smooth(data = step_df1, method = "lm", aes(y = predicted, color = "predicted")) +
  geom_smooth(data = step_df2, method = "lm", aes(y = predicted, color = "predicted")) +
  geom_smooth(data = step_df3, method = "lm", aes(y = predicted, color = "predicted")) +
  scale_colour_manual("", values = c("actual" = "blue",
                                     "predicted" = "red")) +
  xlab("income06") + ylab("egalit_scale") +
  labs(title = "Predicted egalit_scale",
       subtitle = "Step function with 3 cuts")

```



## Interpret the results

Like the polynomial regression model above, the step function (with 3 cuts) is of middling accuracy, judging from the R-squared values of the model and the fit of the model to the data.

### 3. NATURAL REGRESSION SPLINE

Fit a natural regression spline to predict `egalit_scale` as a function of `income06`

Use 10-fold cross-validation to select optimal number of degrees of freedom

```
degree <- 1:10
natreg_rsquared <- rep(0,10)
natreg_RMSE <- rep(0,10)

for (i in 1:10) {
  natreg_formula <- bquote(egalit_scale ~ ns(income06, df = .(i)))
  natreg_mod <- train(as.formula(natreg_formula),
                     data = train,
                     method = "lm",
                     trControl = train_control)
  natreg_rsquared[i] <- natreg_mod$results$Rsquared
  natreg_RMSE[i] <- natreg_mod$results$RMSE
}

natreg_table <- cbind(degree, natreg_rsquared, natreg_RMSE) %>%
  as.data.frame() %>%
  arrange(natreg_RMSE)
natreg_table
```

##	degree	natreg_rsquared	natreg_RMSE
## 1	3	0.06335959	9.325581
## 2	2	0.06784346	9.337192
## 3	10	0.06219866	9.337570
## 4	7	0.06458164	9.343013
## 5	1	0.06037112	9.343996
## 6	9	0.06189211	9.350200
## 7	8	0.05890278	9.351720
## 8	5	0.05830273	9.352151
## 9	6	0.06068609	9.353680
## 10	4	0.05981056	9.353758

Select 3 as number of degrees of freedom, given that it has the lowest RMSE (accuracy) and is a relatively low number of degrees (interpretability).

```
natreg_mod <- lm(egalit_scale ~ ns(income06, df = 3), data = train)
natreg_pred <- predict(step_mod, newdata = test)

natreg_df <- data.frame(income = test$income06,
                       actual = test$egalit_scale,
                       predicted = natreg_pred)
```

Present the results of the optimal model

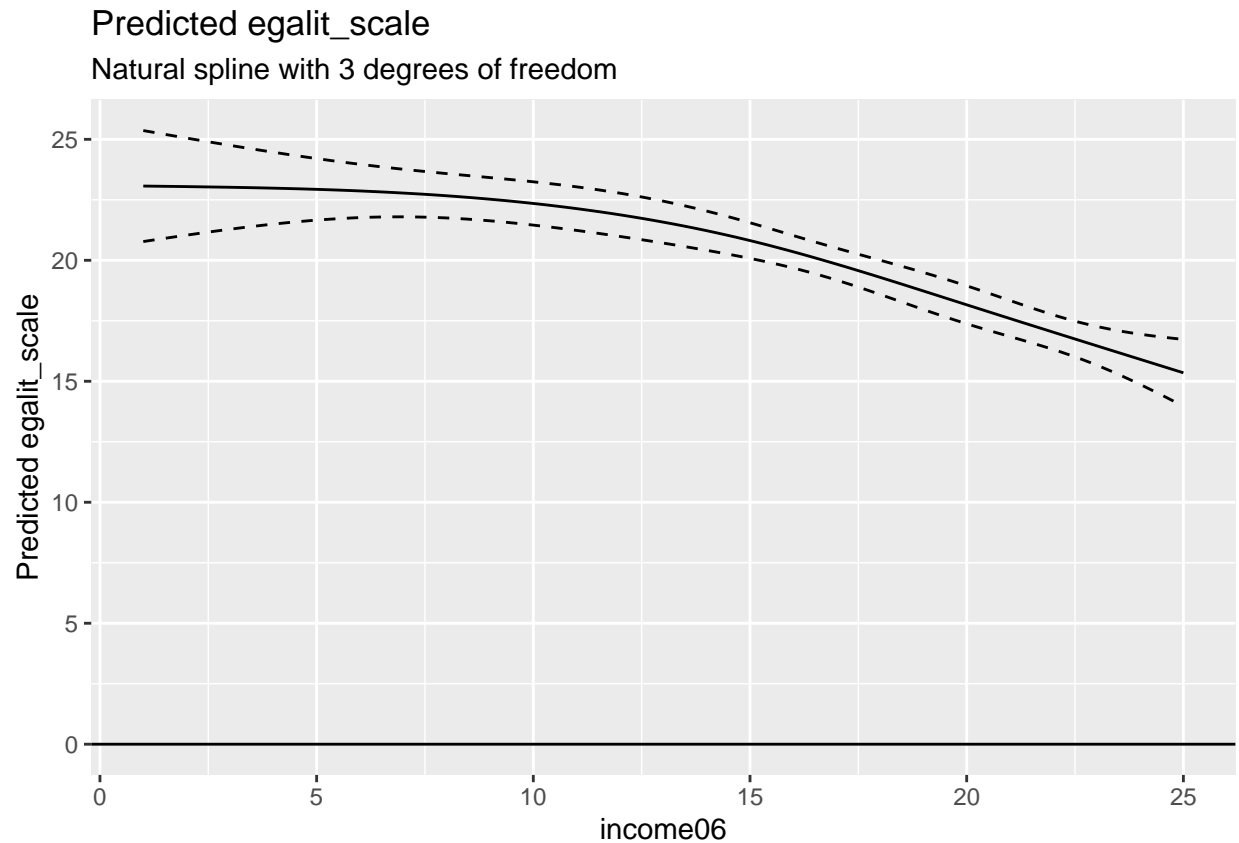
```

cplot(natreg_mod, "income06", what = "prediction", n = 101, draw = FALSE) %>%
  ggplot(aes(x = xvals)) +
  geom_line(aes(y = yvals)) +
  geom_line(aes(y = upper), linetype = 2) +
  geom_line(aes(y = lower), linetype = 2) +
  geom_hline(yintercept = 0, linetype = 1) +
  labs(title = "Predicted egalit_scale",
        subtitle = "Natural spline with 3 degrees of freedom",
        x = "income06",
        y = "Predicted egalit_scale")

```

##	xvals	yvals	upper	lower
## 1	1.00	23.06899	25.36250	20.77548
## 2	1.24	23.06347	25.28680	20.84015
## 3	1.48	23.05790	25.21163	20.90417
## 4	1.72	23.05221	25.13704	20.96738
## 5	1.96	23.04635	25.06307	21.02962
## 6	2.20	23.04025	24.98976	21.09073
## 7	2.44	23.03386	24.91716	21.15055
## 8	2.68	23.02711	24.84533	21.20890
## 9	2.92	23.01996	24.77432	21.26561
## 10	3.16	23.01235	24.70420	21.32049
## 11	3.40	23.00420	24.63502	21.37338
## 12	3.64	22.99547	24.56687	21.42408
## 13	3.88	22.98610	24.49981	21.47240
## 14	4.12	22.97603	24.43392	21.51815
## 15	4.36	22.96520	24.36929	21.56112
## 16	4.60	22.95356	24.30599	21.60112
## 17	4.84	22.94103	24.24411	21.63795
## 18	5.08	22.92757	24.18374	21.67140
## 19	5.32	22.91312	24.12495	21.70129
## 20	5.56	22.89762	24.06782	21.72742





The natural spline model (with 3 degrees of freedom) shows that there is an inverse relationship between `income06` and `egalit_scale`. The higher an individual's income, the less egalitarian they are likely to be. As such, the natural spline model supports the general findings from the earlier polynomial regression model.

#### 4(a). LINEAR REGRESSION

Perform appropriate data pre-processing

Standardize numerical features

```
test <- test %>%  
  mutate_if(is.integer, standardize)  
  
train <- train %>%  
  mutate_if(is.integer, standardize)
```

Estimate model using all available predictors

Use 10-fold cross-validation to estimate model's performance using MSE

```

train_control <- trainControl(method = "CV", number = 10)
linear_mod <- train(egalit_scale ~ .,
  data = train,
  method = "lm",
  trControl = train_control)
linear_pred <- linear_mod %>% predict(test)
linear_mse <- (RMSE(linear_pred, test$egalit_scale))^2
linear_mse

```

```
## [1] 0.6954183
```

## 4(b). ELASTIC NET REGRESSION

Perform appropriate hyperparameter tuning

Estimate model using all available predictors

Use 10-fold cross-validation to estimate model's performance using MSE

```

enet_mod <- train(egalit_scale ~ .,
  data = train,
  method = "glmnet",
  trControl = train_control,
  tuneGrid = expand.grid(alpha = seq(0, 1, 0.1),
    lambda = seq(0.001, 0.1, 0.001)))

enet_pred <- enet_mod %>% predict(test)
enet_mse <- (RMSE(enet_pred, test$egalit_scale))^2
enet_mse

```

```
## [1] 0.6775607
```

## 4(c). PRINCIPAL COMPONENT REGRESSION

Perform appropriate hyperparameter tuning

Estimate model using all available predictors

```

pcr_mod <- train(egalit_scale ~ .,
  data = train,
  method = "pcr",
  trControl = train_control,
  tuneLength = 10)

```

Use 10-fold cross-validation to estimate model's performance using MSE

```
pcr_pred <- pcr_mod %>% predict(test)
pcr_mse <- (RMSE(pcr_pred, test$egalit_scale))^2
pcr_mse
```

```
## [1] 0.8520385
```

## 4(d). PARTIAL LEAST SQUARES REGRESSION

Perform hyperparameter tuning

Estimate model using all available predictors

```
pls_mod <- train(egalit_scale ~ .,
  data = train,
  method = "pls",
  trControl = train_control,
  tuneLength = 10)
```

Use 10-fold cross-validation to estimate model's performance using MSE

```
pls_pred <- pls_mod %>% predict(test)
pls_mse <- (RMSE(pls_pred, test$egalit_scale))^2
pls_mse
```

```
## [1] 0.6866766
```

## 5. FEATURE INTERACTION PLOTS

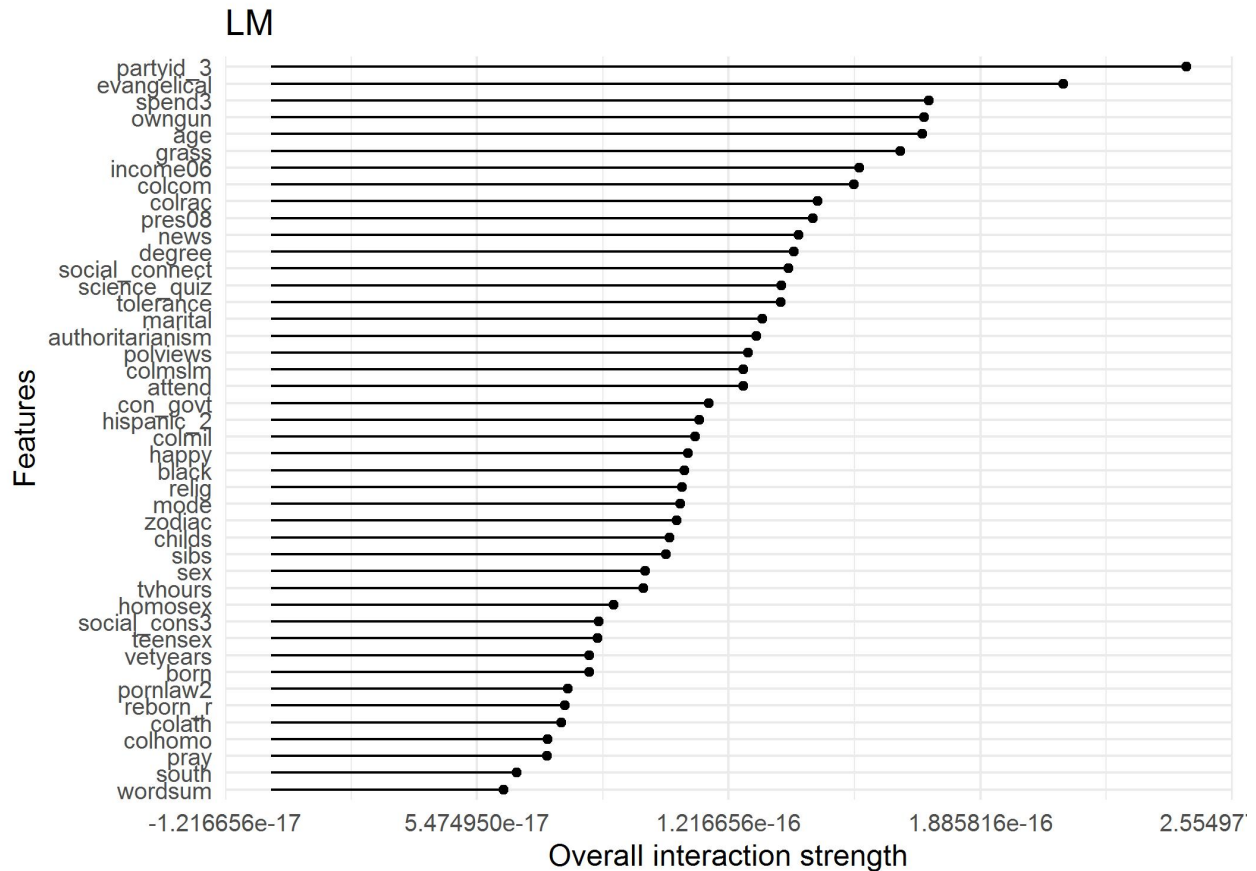
Linear Regression

Plot

```
predictor_lm <- Predictor$new(
  model = linear_mod,
  data = features,
  y = response,
  predict.fun = pred,
  class = "classification"
)
```

```
interact_lm <- Interaction$new(predictor_lm)
```

```
plot(interact_lm) +  
  ggtitle("LM") +  
  theme_minimal(base_size = 12)
```



## Discussion

In this linear model, `party_id3`, `evangelical`, and `spend3` are the three features with the highest interaction strength. `partyid3` is by far the strongest, followed by `evangelical`. `spend3` is not much stronger than the variables after it.

## Elastic Net Regression

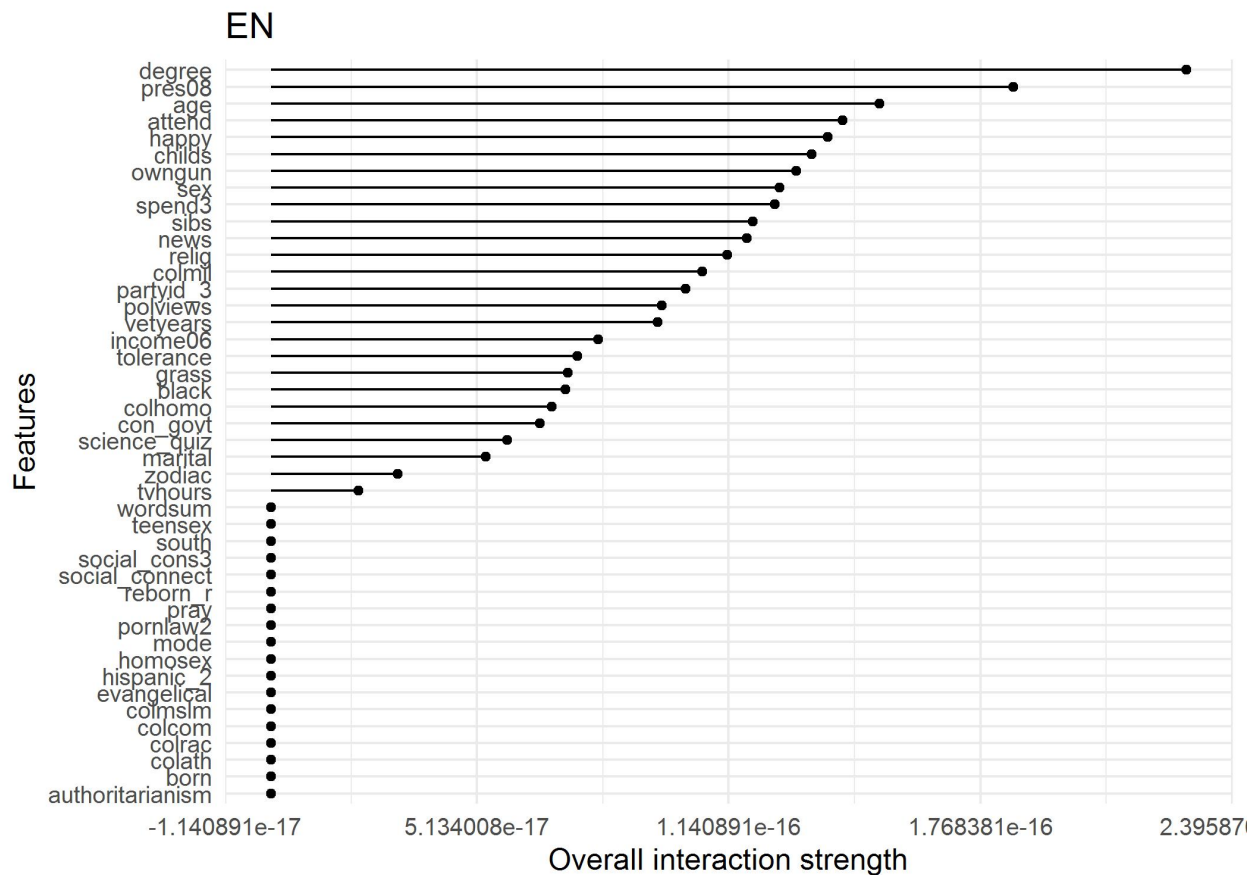
### Plot

```
predictor_enet <- Predictor$new(  
  model =enet_mod,  
  data = features,  
  y = response,  
  predict.fun = pred,
```

```
class = "classification"
)
```

```
interact_enet <- Interaction$new(predictor_enet)
```

```
plot(interact_enet) +
  ggtitle("EN") +
  theme_minimal(base_size = 12)
```



### Discussion In this elastic net model, **degree**, **pres08**, and **age** are the three features with the highest interaction strength. **degree** is by far the strongest, followed by **pres08**. **age08** is only marginally stronger than the variables after it. Many variables are of extremely low interaction strength, from **wordsum** to **authoritarianism**.

## Principal Component Regression


### Plot

```
predictor_pcr <- Predictor$new(
  model = pcr_mod,
  data = features,
  y = response,
  predict.fun = pred,
```

```
class = "classification"
)
```

```
interact_pcr <- Interaction$new(predictor_pcr)
```

```
plot(interact_pcr) +
  ggtitle("PCR") +
  theme_minimal(base_size = 12)
```



## Discussion

In this PCR model, `income06`, `social_connect`, and `attend` are the three features with the highest interaction strength. The three features are clustered together, with interaction strengths that are significantly higher than the next two features (`born` and `con_govt`).

## Partial Least Squares Regression

### Plot

```
predictor_pls <- Predictor$new(
  model = pls_mod,
  data = features,
```

```

y = response,
predict.fun = pred,
class = "classification"
)

```

```

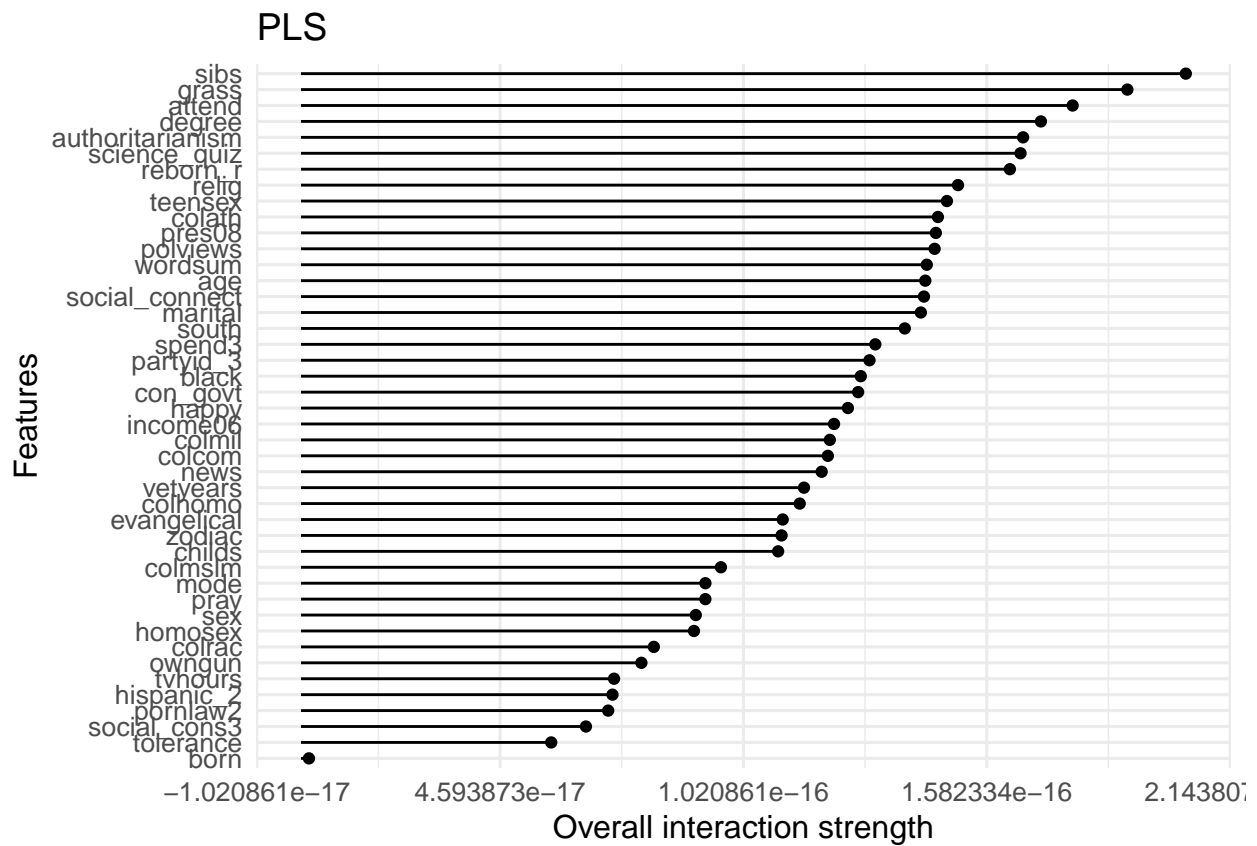
interact_pls <- Interaction$new(predictor_pls)

```

```

plot(interact_pls) +
  ggtitle("PLS") +
  theme_minimal(base_size = 12)

```



## Discussion

In this PLS model, **sibs**, **grass**, and **attend** are the three features with the highest interaction strength. **sibs** is by far the strongest. **grass** and **attend** follow, decreasing in strength by about the same quantities.

## Comparison

Observe that the top three features of each model are largely different:

- \* LM: partyid\_3, evangelical, spend3
- \* EN: degree, pres08, age
- \* PCR: income06, social\_connect, attend
- \* PLS: sibs, grass, attend

The only repeated feature is **attend**.

Observe also that the clustering of feature interaction strengths differs between models:

- \* LM: 1 feature that is clearly strongest, 1 feature that is clearly the next-strongest, followed by many other features of noticeable strength

- \* EN: 1 feature that is clearly strongest, 1 feature that is clearly the next-strongest, followed by several other features of noticeable strength, and several other features of negligible strength

- \* PCR: 3 features that are clearly strongest, followed by small clusters of features of decreasing strength

- \* PLS: 1 feature that is clearly strongest, 1 feature that is clearly the next-strongest, 1 feature that is the next-strongest after, and large clusters of features of decreasing strength, ending with 1 feature of negligible strength

The differences between the models may suggest that some models are better suited to the GSS data than others, or even that the models examined may not capture the actual features that weigh the most heavily on the response. In either cases, different models carry different assumptions about the data, which influence how the models identify the most important features.